

Implementation Security In Cryptography

Lecture 11: Entering the World of Attacks

Recap

- In the last lecture
 - Compact design of AES

Today

- Few more words about implementations
- Entering the world of attacks.

AES Once Again

- Well, we have seen how it is done in hardware..
- But what about software?
- A popular, extremely fast, yet terrible approach — T tables
 - Used quite a lot in OpenSSL
 - Not used anymore due to several **cache timing attacks**
 - But table based secure implementations also do exist

AES T Tables

$$\begin{pmatrix} 2 & 1 & 1 & 3 \\ 3 & 2 & 1 & 1 \\ 1 & 3 & 2 & 1 \\ 1 & 1 & 3 & 2 \end{pmatrix} \begin{pmatrix} c_0 \\ c_1 \\ c_2 \\ c_3 \end{pmatrix} = \begin{pmatrix} s_0 \\ s_1 \\ s_2 \\ s_3 \end{pmatrix}$$

- Let us consider the MixColumns operation
- In software, each 32-bit AES column is a uint32 variable.

$$s_0 = 2c_0 \oplus 3c_1 \oplus 1c_2 \oplus 1c_3$$

$$s_1 = 1c_0 \oplus 2c_1 \oplus 3c_2 \oplus 1c_3$$

$$s_2 = 1c_0 \oplus 1c_1 \oplus 2c_2 \oplus 3c_3$$

$$s_3 = 3c_0 \oplus 1c_1 \oplus 1c_2 \oplus 2c_3$$

AES T Tables

$$\begin{pmatrix} 2 & 1 & 1 & 3 \\ 3 & 2 & 1 & 1 \\ 1 & 3 & 2 & 1 \\ 1 & 1 & 3 & 2 \end{pmatrix} \begin{pmatrix} c_0 \\ c_1 \\ c_2 \\ c_3 \end{pmatrix} = \begin{pmatrix} s_0 \\ s_1 \\ s_2 \\ s_3 \end{pmatrix}$$

- Let us consider the MixColumns operation
- In software, each 32-bit AES column is a uint32 variable.
- So we can store s in a uint32 variable.
- Here each s_i is 8-bit, so we can simply do concatenation to get 32-bits.

$$\begin{aligned} s_0 &= 2c_0 \oplus 3c_1 \oplus 1c_2 \oplus 1c_3 \\ s_1 &= 1c_0 \oplus 2c_1 \oplus 3c_2 \oplus 1c_3 \\ s_2 &= 1c_0 \oplus 1c_1 \oplus 2c_2 \oplus 3c_3 \\ s_3 &= 3c_0 \oplus 1c_1 \oplus 1c_2 \oplus 2c_3 \end{aligned}$$

```
s = s0 | s1 | s2 | s3
s = 2*c0 ^ 3*c1 ^ 1*c2 ^ 1*c3 |
    1*c0 ^ 2*c1 ^ 3*c2 ^ 1*c3 |
    1*c0 ^ 1*c1 ^ 2*c2 ^ 3*c3 |
    3*c0 ^ 1*c1 ^ 1*c2 ^ 2*c3
```

AES T Tables

$$\begin{pmatrix} 2 & 1 & 1 & 3 \\ 3 & 2 & 1 & 1 \\ 1 & 3 & 2 & 1 \\ 1 & 1 & 3 & 2 \end{pmatrix} \begin{pmatrix} c_0 \\ c_1 \\ c_2 \\ c_3 \end{pmatrix} = \begin{pmatrix} s_0 \\ s_1 \\ s_2 \\ s_3 \end{pmatrix}$$

- Here each s_i is 8-bit, so we can simply do concatenation to get 32-bits.
- Now we can also rearrange the terms.
 - Why this is beneficial??

```
s = s0 | s1 | s2 | s3  
s = 2*c0 ^ 3*c1 ^ 1*c2 ^ 1*c3 |  
    1*c0 ^ 2*c1 ^ 3*c2 ^ 1*c3 |  
    1*c0 ^ 1*c1 ^ 2*c2 ^ 3*c3 |  
    3*c0 ^ 1*c1 ^ 1*c2 ^ 2*c3
```

```
s = (2*c0 | 1*c0 | 1*c0 | 3*c0) ^  
(3*c1 | 2*c1 | 1*c1 | 1*c1) ^  
(1*c2 | 3*c2 | 2*c2 | 1*c2) ^  
(1*c3 | 1*c3 | 3*c3 | 2*c3)
```

AES T Tables

$$\begin{pmatrix} 2 & 1 & 1 & 3 \\ 3 & 2 & 1 & 1 \\ 1 & 3 & 2 & 1 \\ 1 & 1 & 3 & 2 \end{pmatrix} \begin{pmatrix} c_0 \\ c_1 \\ c_2 \\ c_3 \end{pmatrix} = \begin{pmatrix} s_0 \\ s_1 \\ s_2 \\ s_3 \end{pmatrix}$$

- Here each s_i is 8-bit, so we can simply do concatenation to get 32-bits.
- Now we can also rearrange the terms.
 - Why this is beneficial??
 - Observe that you can compute each term being XORed only from a c_i

```
s = s0 | s1 | s2 | s3  
s = 2*c0 ^ 3*c1 ^ 1*c2 ^ 1*c3 |  
    1*c0 ^ 2*c1 ^ 3*c2 ^ 1*c3 |  
    1*c0 ^ 1*c1 ^ 2*c2 ^ 3*c3 |  
    3*c0 ^ 1*c1 ^ 1*c2 ^ 2*c3
```

```
s = (2*c0 | 1*c0 | 1*c0 | 3*c0) ^  
(3*c1 | 2*c1 | 1*c1 | 1*c1) ^  
(1*c2 | 3*c2 | 2*c2 | 1*c2) ^  
(1*c3 | 1*c3 | 3*c3 | 2*c3)
```

AES T Tables

$$\begin{pmatrix} 2 & 1 & 1 & 3 \\ 3 & 2 & 1 & 1 \\ 1 & 3 & 2 & 1 \\ 1 & 1 & 3 & 2 \end{pmatrix} \begin{pmatrix} c_0 \\ c_1 \\ c_2 \\ c_3 \end{pmatrix} = \begin{pmatrix} s_0 \\ s_1 \\ s_2 \\ s_3 \end{pmatrix}$$

- Why this is beneficial??
 - Observe that you can compute each term being XORed only from a c_i
 - Since c_i is 8-bit, so we can compute all possible 256 values and store them in. A table.
 - Same can be done for all c_i
 - One table for each
 - **Catch:** Table lookup is much faster than a finite field operation.

```
s = s0 | s1 | s2 | s3  
s = 2*c0 ^ 3*c1 ^ 1*c2 ^ 1*c3 |  
    1*c0 ^ 2*c1 ^ 3*c2 ^ 1*c3 |  
    1*c0 ^ 1*c1 ^ 2*c2 ^ 3*c3 |  
    3*c0 ^ 1*c1 ^ 1*c2 ^ 2*c3
```

```
s = (2*c0 | 1*c0 | 1*c0 | 3*c0) ^  
(3*c1 | 2*c1 | 1*c1 | 1*c1) ^  
(1*c2 | 3*c2 | 2*c2 | 1*c2) ^  
(1*c3 | 1*c3 | 3*c3 | 2*c3)
```

AES T Tables

$$\begin{pmatrix} 2 & 1 & 1 & 3 \\ 3 & 2 & 1 & 1 \\ 1 & 3 & 2 & 1 \\ 1 & 1 & 3 & 2 \end{pmatrix} \begin{pmatrix} c_0 \\ c_1 \\ c_2 \\ c_3 \end{pmatrix} = \begin{pmatrix} s_0 \\ s_1 \\ s_2 \\ s_3 \end{pmatrix}$$

- Why this is beneficial??

- We denote these tables as $te0$, $te1$, $te2$, and $te3$.
- The operation is as follows:

```
te0[i] = 2*i | 1*i | 1*i | 3*i  
te1[i] = 3*i | 2*i | 1*i | 1*i  
te2[i] = 1*i | 3*i | 2*i | 1*i  
te3[i] = 1*i | 1*i | 3*i | 2*i
```

```
s = te0[c0] ^ te1[c1] ^ te2[c2] ^ te3[c3]
```

```
s = s0 | s1 | s2 | s3  
s = 2*c0 ^ 3*c1 ^ 1*c2 ^ 1*c3 |  
    1*c0 ^ 2*c1 ^ 3*c2 ^ 1*c3 |  
    1*c0 ^ 1*c1 ^ 2*c2 ^ 3*c3 |  
    3*c0 ^ 1*c1 ^ 1*c2 ^ 2*c3
```

```
s = (2*c0 | 1*c0 | 1*c0 | 3*c0) ^  
(3*c1 | 2*c1 | 1*c1 | 1*c1) ^  
(1*c2 | 3*c2 | 2*c2 | 1*c2) ^  
(1*c3 | 1*c3 | 3*c3 | 2*c3)
```

It does not ends here...

$$\begin{pmatrix} 2 & 1 & 1 & 3 \\ 3 & 2 & 1 & 1 \\ 1 & 3 & 2 & 1 \\ 1 & 1 & 3 & 2 \end{pmatrix} \begin{pmatrix} c_0 \\ c_1 \\ c_2 \\ c_3 \end{pmatrix} = \begin{pmatrix} s_0 \\ s_1 \\ s_2 \\ s_3 \end{pmatrix}$$

- Can we do better than this?

- ```
te0[i] = 2*i | 1*i | 1*i | 3*i
te1[i] = 3*i | 2*i | 1*i | 1*i
te2[i] = 1*i | 3*i | 2*i | 1*i
te3[i] = 1*i | 1*i | 3*i | 2*i
```

```
s = te0[c0] ^ te1[c1] ^ te2[c2] ^ te3[c3]
```

```
s = s0 | s1 | s2 | s3
s = 2*c0 ^ 3*c1 ^ 1*c2 ^ 1*c3 |
 1*c0 ^ 2*c1 ^ 3*c2 ^ 1*c3 |
 1*c0 ^ 1*c1 ^ 2*c2 ^ 3*c3 |
 3*c0 ^ 1*c1 ^ 1*c2 ^ 2*c3
```

```
s = (2*c0 | 1*c0 | 1*c0 | 3*c0) ^
 (3*c1 | 2*c1 | 1*c1 | 1*c1) ^
 (1*c2 | 3*c2 | 2*c2 | 1*c2) ^
 (1*c3 | 1*c3 | 3*c3 | 2*c3)
```

# It does not ends here...

```
s = (2*c0 | 1*c0 | 1*c0 | 3*c0) ^
 (3*c1 | 2*c1 | 1*c1 | 1*c1) ^
 (1*c2 | 3*c2 | 2*c2 | 1*c2) ^
 (1*c3 | 1*c3 | 3*c3 | 2*c3)
```

- Can we do better than this?
- Observe that:  $c_0|c_1|c_2|c_3$  can be
  - $S[b_0]|S[b_5]|S[b_{10}]|S[b_{15}]$  Or,  $S[b_4]|S[b_9]|S[b_{14}]|S[b_3]$  Or  
 $S[b_8]|S[b_{13}]|S[b_2]|S[b_7]$     $S[b_{12}]|S[b_1]|S[b_6]|S[b_{11}]$
  - So we can merge subtypes and shift rows in a table

```
R0 = (2*S[b0] | 1*S[b0] | 1*S[b0] | 3*S[b0]) ^
 (3*S[b5] | 2*S[b5] | 1*S[b5] | 1*S[b5]) ^
 (1*S[b10] | 3*S[b10] | 2*S[b10] | 1*S[b10]) ^
 (1*S[b15] | 1*S[b15] | 3*S[b15] | 2*S[b15]) ^

R1 = (2*S[b4] | 1*S[b4] | 1*S[b4] | 3*S[b4]) ^
 (3*S[b9] | 2*S[b9] | 1*S[b9] | 1*S[b9]) ^
 (1*S[b14] | 3*S[b14] | 2*S[b14] | 1*S[b14]) ^
 (1*S[b3] | 1*S[b3] | 3*S[b3] | 2*S[b3]) ^

R2 = (2*S[b8] | 1*S[b8] | 1*S[b8] | 3*S[b8]) ^
 (3*S[b13] | 2*S[b13] | 1*S[b13] | 1*S[b13]) ^
 (1*S[b2] | 3*S[b2] | 2*S[b2] | 1*S[b2]) ^
 (1*S[b7] | 1*S[b7] | 3*S[b7] | 2*S[b7]) ^

R3 = (2*S[b12] | 1*S[b12] | 1*S[b12] | 3*S[b12]) ^
 (3*S[b1] | 2*S[b1] | 1*S[b1] | 1*S[b1]) ^
 (1*S[b6] | 3*S[b6] | 2*S[b6] | 1*S[b6]) ^
 (1*S[b11] | 1*S[b11] | 3*S[b11] | 2*S[b11]) ^
```

# It does not ends here...

- So finally

```
R0 = te0[b0] ^ te1[b5] ^ te2[b10] ^ te3[b15]
R1 = te0[b4] ^ te1[b9] ^ te2[b14] ^ te3[b3]
R2 = te0[b8] ^ te1[b13] ^ te2[b2] ^ te3[b7]
R3 = te0[b12] ^ te1[b1] ^ te2[b6] ^ te3[b11]
```

# It doe

- So finall

```
R0 = te0[b0]
R1 = te0[b4]
R2 = te0[b8]
R3 = te0[b12]
```

```
// Initialize the state, stored in s0, s1, s2 and s3
s0 := uint32(src[0])<<24 | uint32(src[1])<<16 | uint32(src[2])<<8 | uint32(src[3])
s1 := uint32(src[4])<<24 | uint32(src[5])<<16 | uint32(src[6])<<8 | uint32(src[7])
s2 := uint32(src[8])<<24 | uint32(src[9])<<16 | uint32(src[10])<<8 | uint32(src[11])
s3 := uint32(src[12])<<24 | uint32(src[13])<<16 | uint32(src[14])<<8 | uint32(src[15])

R0 = te0[b0]
R1 = te0[b4] // Add the first round key to the state
R2 = te0[b8]
R3 = te0[b12]
s0 ^= xk[0]
s1 ^= xk[1]
s2 ^= xk[2]
s3 ^= xk[3]

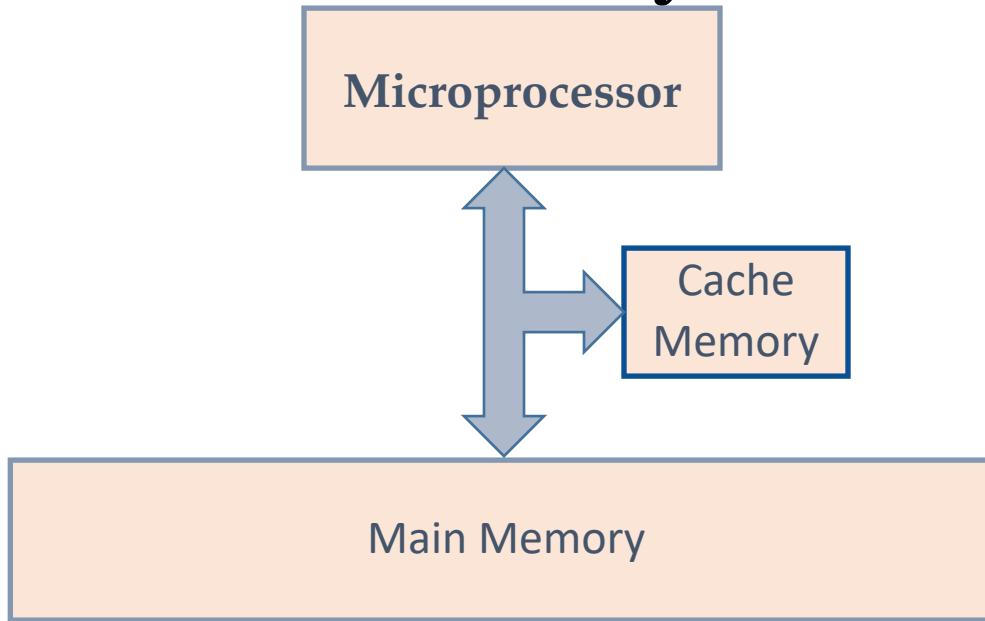
for i:= 1; i < nr; i++ {
 // This performs SubBytes + ShiftRows + MixColumns + AddRoundKey
 tmp0 = te0[s0>>24] ^ te1[s1>>16&0xff] ^ te2[s2>>8&0xff] ^ te3[s3&0xff] ^ xk[4*i]
 tmp1 = te0[s1>>24] ^ te1[s2>>16&0xff] ^ te2[s3>>8&0xff] ^ te3[s0&0xff] ^ xk[4*i+1]
 tmp2 = te0[s2>>24] ^ te1[s3>>16&0xff] ^ te2[s0>>8&0xff] ^ te3[s1&0xff] ^ xk[4*i+2]
 tmp3 = te0[s3>>24] ^ te1[s0>>16&0xff] ^ te2[s1>>8&0xff] ^ te3[s2&0xff] ^ xk[4*i+3]

 s0, s1, s2, s3 = tmp0, tmp1, tmp2, tmp3
}
```

# But As We Said...

- The tables are stored in cache memory of your system.
- AES accesses this table depending on the secret key values
- An adversary, who is able to measure the time for each encryption operation, and also using the same cache can do something so that it can recover the secret key!!!
- Such attacks are called cache timing attacks...

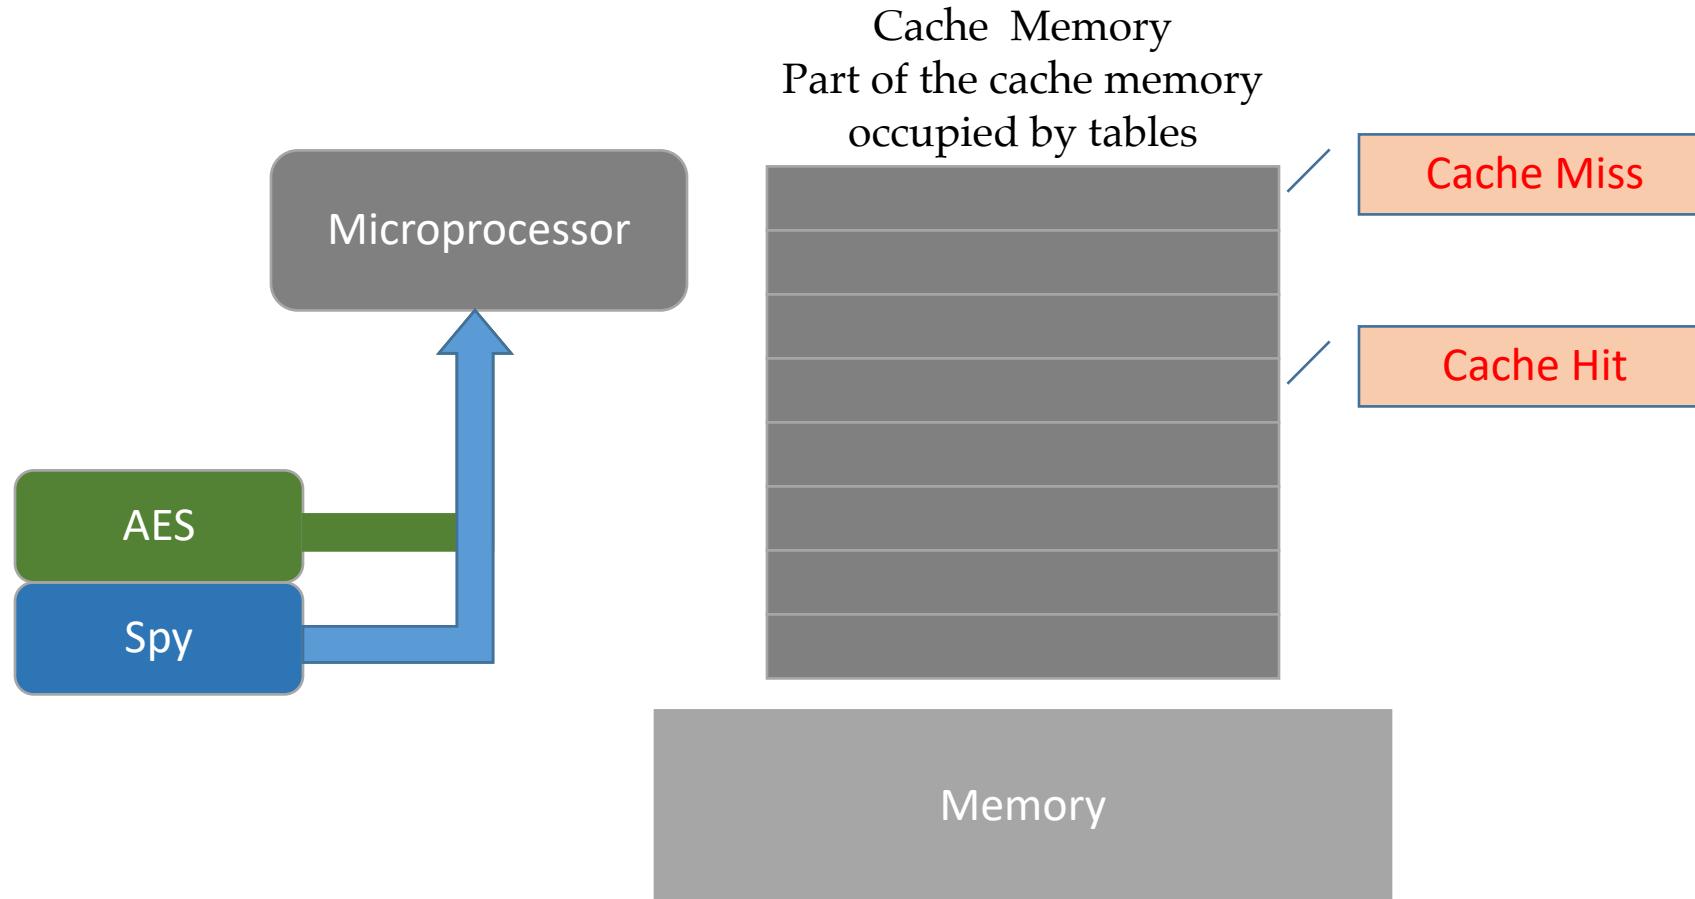
# Attacks due to Memory Wall



- If there is a *Cache Hit*
  - Access time is less
  - Power Consumption is less
- If there is a *Cache Miss*
  - Access time is more
  - Power Consumption is more

# Timing Attacks due to Cache Memory

- Uses a spy program to determine cache behavior



# Bitslicing

- Simply speaking, implement the software like hardware
- Results in constant-time crypto
- Idea: Let's say you are running on a 32-bit machine.
  - 32-bit registers
  - Logical AND, OR, NOT, XOR.
  - Also consider your block cipher in terms of these gates.

# Bitslicing: A Simple Example

- Let us consider the first equation only.

- It can be written as follows:

and  $t_1, x_1, x_4$

and  $t_2, t_1, x_2$

and  $t_3, t_1, x_3$

xor  $t_4, t_2, t_3$

And so on...

$$y_1 = x_1 x_2 x_4 + x_1 x_3 x_4 + x_1 + x_2 x_3 x_4 + x_2 x_3 + x_3 + x_4 + 1$$

$$y_2 = x_1 x_2 x_4 + x_1 x_3 x_4 + x_1 x_3 + x_1 x_4 + x_1 + x_2 + x_3 x_4 + 1$$

$$y_3 = x_1 x_2 x_4 + x_1 x_2 + x_1 x_3 x_4 + x_1 x_3 + x_1 + x_2 x_3 x_4 + x_3$$

$$y_4 = x_1 + x_2 x_3 + x_2 + x_4$$

# Bitslicing: A Simple Example

- Let us consider the first equation only.

- It can be written as follows:

and  $t_1, x_1, x_4$

and  $t_2, t_1, x_2$

and  $t_3, t_1, x_3$

xor  $t_4, t_2, t_3$

And so on...

$$\begin{aligned}y_1 &= x_1x_2x_4 + x_1x_3x_4 \\&\quad + x_1 + x_2x_3x_4 + x_2x_3 + x_3 + x_4 + 1 \\y_2 &= x_1x_2x_4 + x_1x_3x_4 + x_1x_3 + x_1x_4 + \\&\quad x_1 + x_2 + x_3x_4 + 1 \\y_3 &= x_1x_2x_4 + x_1x_2 + x_1x_3x_4 + x_1x_3 + \\&\quad x_1 + x_2x_3x_4 + x_3 \\y_4 &= x_1 + x_2x_3 + x_2 + x_4\end{aligned}$$

- Now, each of  $t_1, x_1, x_2, \dots$  are mapped to 32 bit registers; but actually they are processing 1-bit values

- So, what to do?

# Bitslicing: A Simple Example

- Let us consider the first equation only.

- It can be written as follows:

and t1, x1, x4

and t2, t1, x2

and t3, t1, x3

xor t4, t2, t3

And so on...

- Pack each register with independent values and process them using the same instruction!!!
    - Easiest case: you can encrypt 32 plaintext together

# Bitslicing: A Simple Example

- Easiest case: you can encrypt 32 plaintext together
- You can parallelize S-Box computations for one plaintext

A 64-bit word

128 words

|                |                |     |               |               |               |            |
|----------------|----------------|-----|---------------|---------------|---------------|------------|
| $b_{63}^0$     | $b_{62}^0$     | ... | $b_3^0$       | $b_2^0$       | $b_1^0$       | $b_0^0$    |
| $b_{127}^0$    | $b_{126}^0$    | ... | $b_{67}^0$    | $b_{66}^0$    | $b_{65}^0$    | $b_{64}^0$ |
| $b_{63}^1$     | $b_{62}^1$     | ... | $b_3^1$       | $b_2^1$       | $b_1^1$       | $b_0^1$    |
| $b_{127}^1$    | $b_{126}^1$    | ... | $b_{67}^1$    | $b_{66}^1$    | $b_{65}^1$    | $b_{64}^1$ |
| ⋮              | ⋮              | ⋮   | ⋮             | ⋮             | ⋮             | ⋮          |
| $b_{63}^{63}$  | $b_{62}^{63}$  | ... | $b_3^{63}$    | $b_2^{63}$    | $b_1^{63}$    | $b_0^{63}$ |
| $b_{127}^{63}$ | $b_{126}^{63}$ | ... | $b_{67}^{63}$ | $b_{66}^{63}$ | $b_{65}^{63}$ | $b_{64}^0$ |

(a) Original storage matrix

A 64-bit bundle

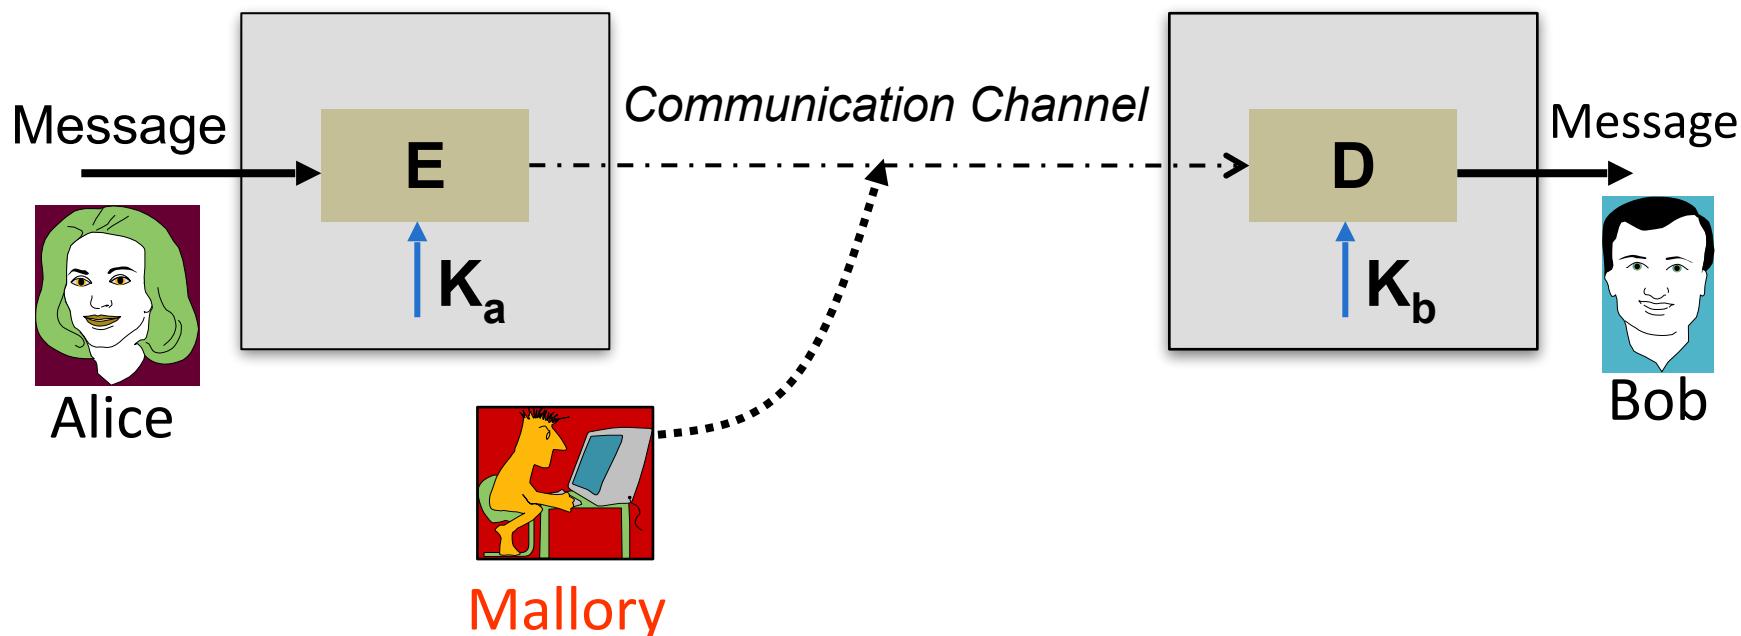
Register 0

|                |                |     |             |             |             |             |
|----------------|----------------|-----|-------------|-------------|-------------|-------------|
| $b_0^{63}$     | $b_0^{62}$     | ... | $b_0^3$     | $b_0^2$     | $b_0^1$     | $b_0^0$     |
| $b_1^{63}$     | $b_1^{62}$     | ... | $b_1^3$     | $b_1^2$     | $b_1^1$     | $b_1^0$     |
| $b_2^{63}$     | $b_2^{62}$     | ... | $b_2^3$     | $b_2^2$     | $b_2^1$     | $b_2^0$     |
| $b_3^{63}$     | $b_3^{62}$     | ... | $b_3^3$     | $b_3^2$     | $b_3^1$     | $b_3^0$     |
| ⋮              | ⋮              | ... | ⋮           | ⋮           | ⋮           | ⋮           |
| $b_{126}^{63}$ | $b_{126}^{62}$ | ... | $b_{126}^3$ | $b_{126}^2$ | $b_{126}^1$ | $b_{126}^0$ |
| $b_{127}^{63}$ | $b_{127}^{62}$ | ... | $b_{127}^3$ | $b_{127}^2$ | $b_{127}^1$ | $b_{127}^0$ |

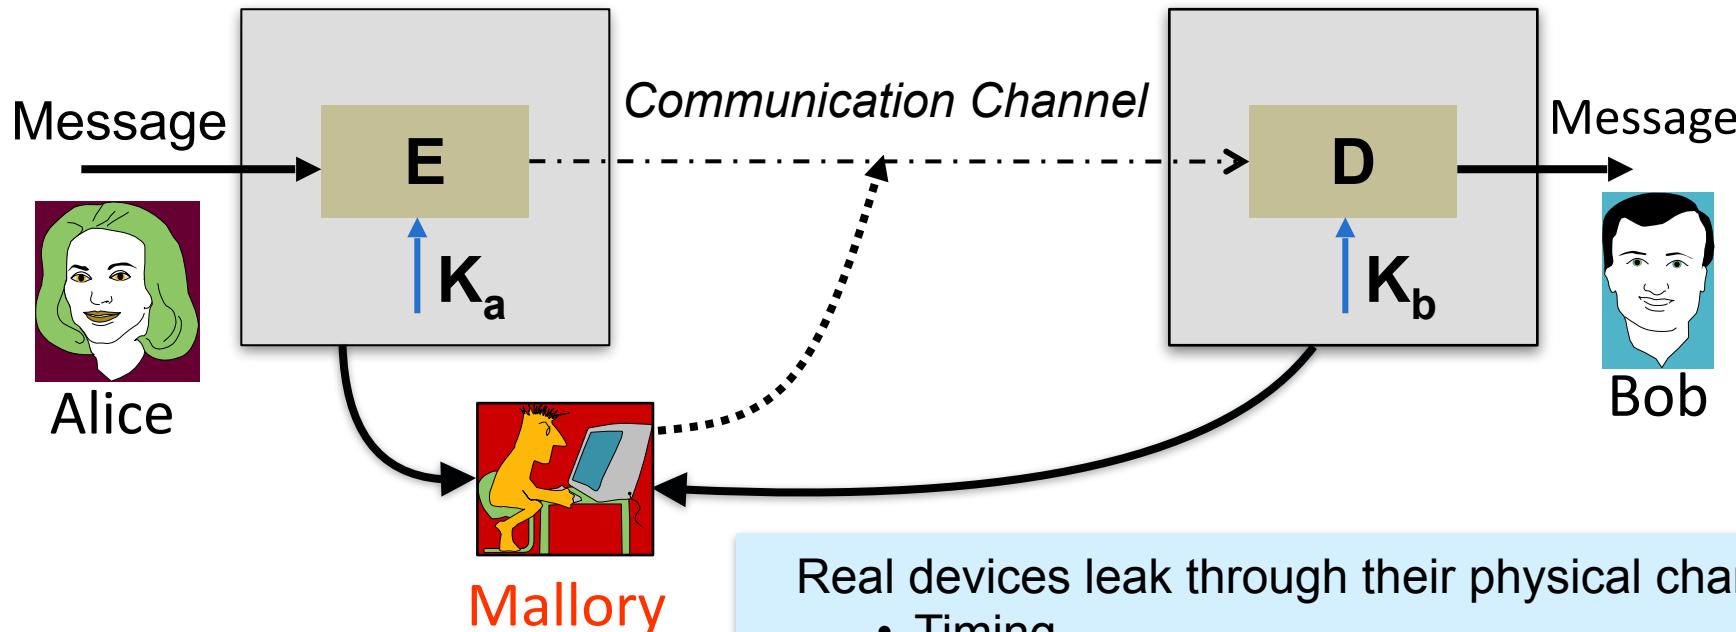
(b) Bit-slice storage matrix

# Side Channel Attacks

# Why do Cryptographers Need Engineers?



# Cryptographic Security: Real World



Strong cryptographic algorithms  
are only the beginning

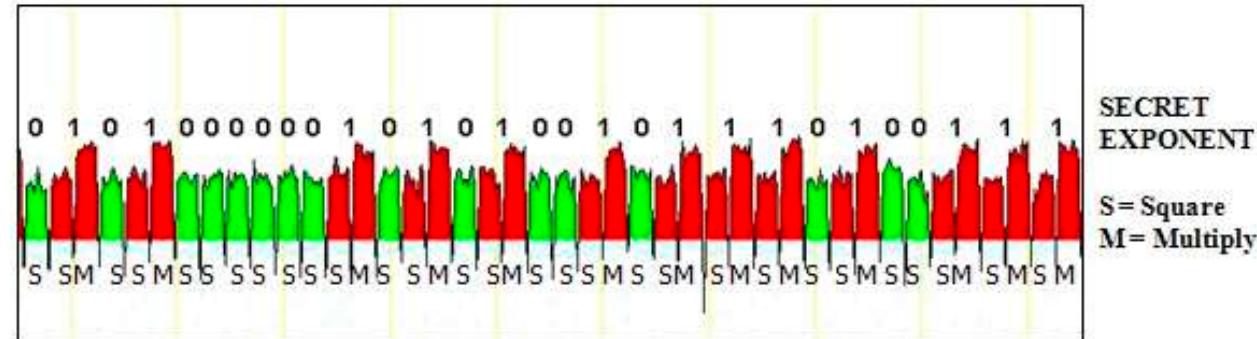
Real devices leak through their physical characteristics

- Timing
- Power consumption
- Electromagnetic Radiation
- Sound
- Faults

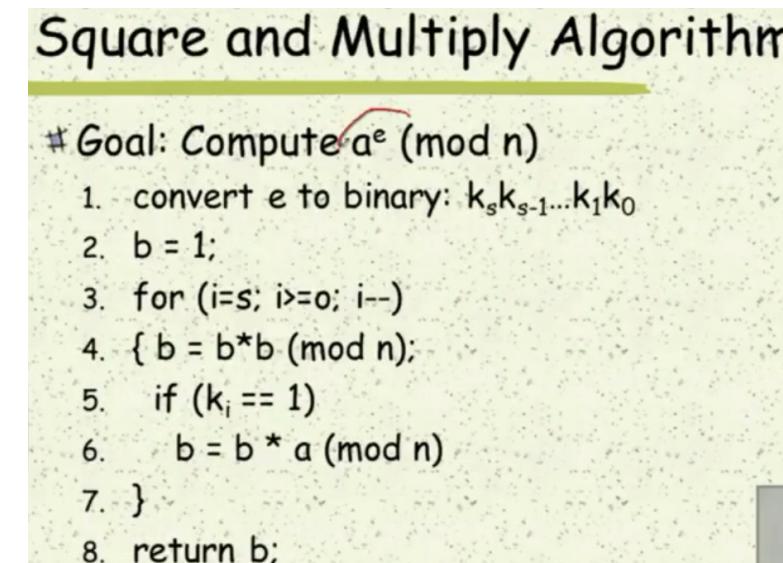
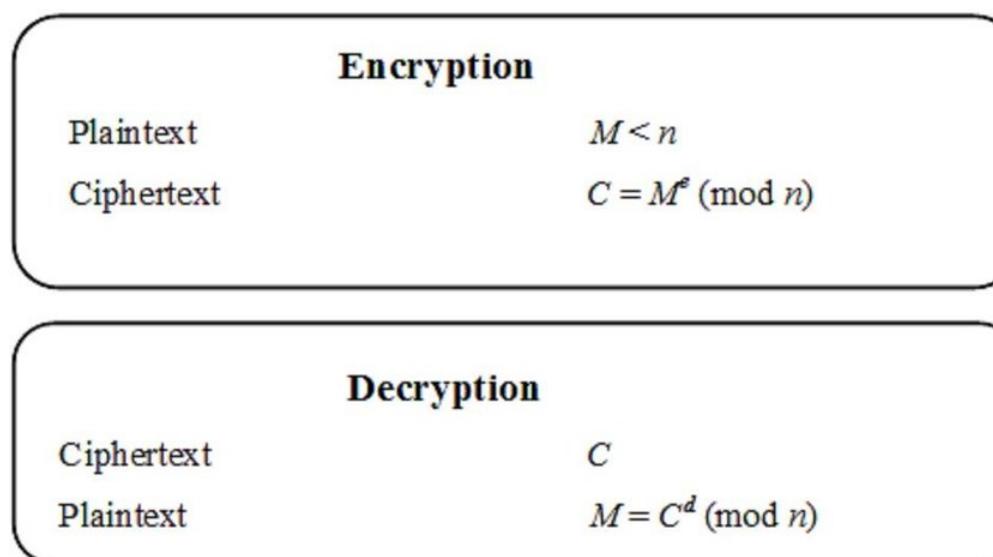
Analysis and mitigation of physical attacks are cryptographic as  
well as engineering problems

# Side-Channel Attacks (SCA)

- The physical channels are correlated with the information being processed
  - Fundamental cause: power consumption is correlated with switching of CMOS transistors ( $0 \rightarrow 1$ ,  $1 \rightarrow 0$ )
    - Typically it is assumed that power consumption is correlated with the Hamming Weight/Distance.
  - If some internal state is exposed, the secret key can be recovered in seconds.

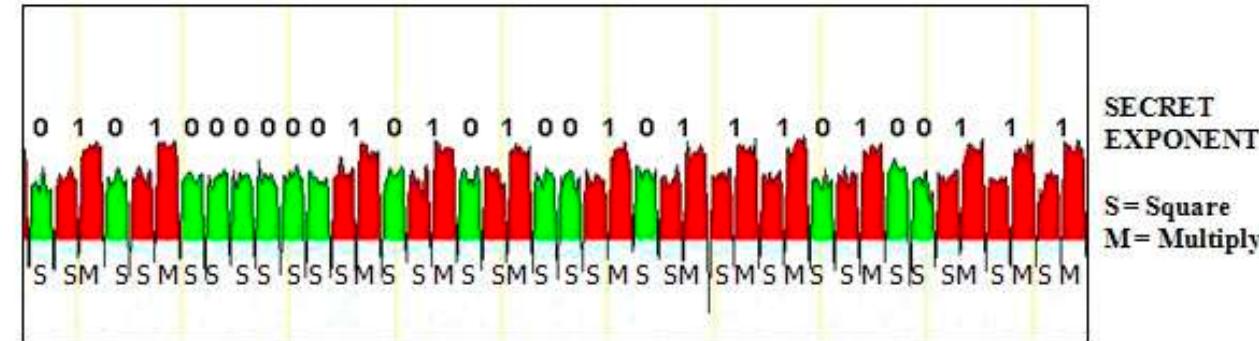
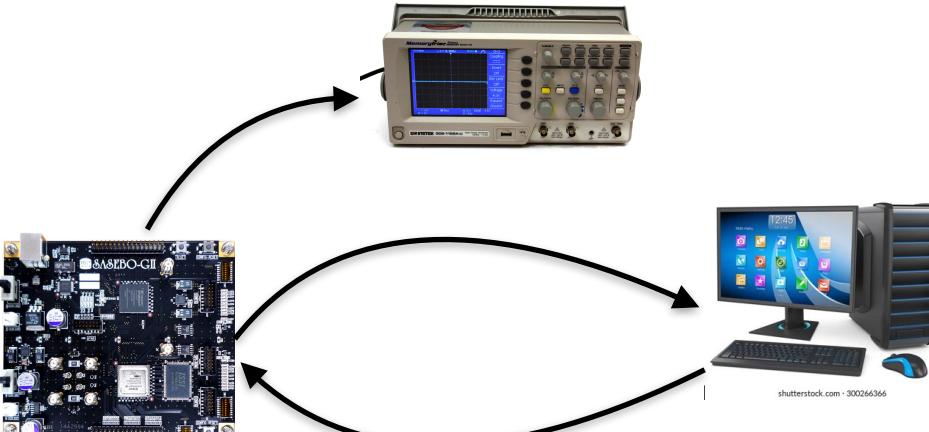


Source: Internet

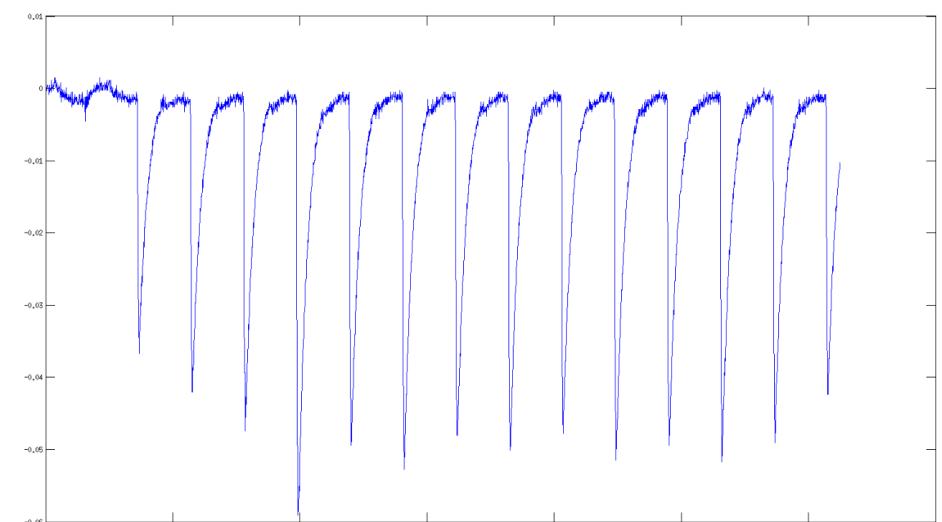


# Side-Channel Attacks (SCA)

- The physical channels are correlated with the information being processed
- Fundamental cause: power consumption is correlated with switching of CMOS transistors ( $0 \rightarrow 1$ ,  $1 \rightarrow 0$ )
  - Typically it is assumed that power consumption is correlated with the Hamming Weight/Distance.
- If some internal state is exposed, the secret key can be recovered in seconds.



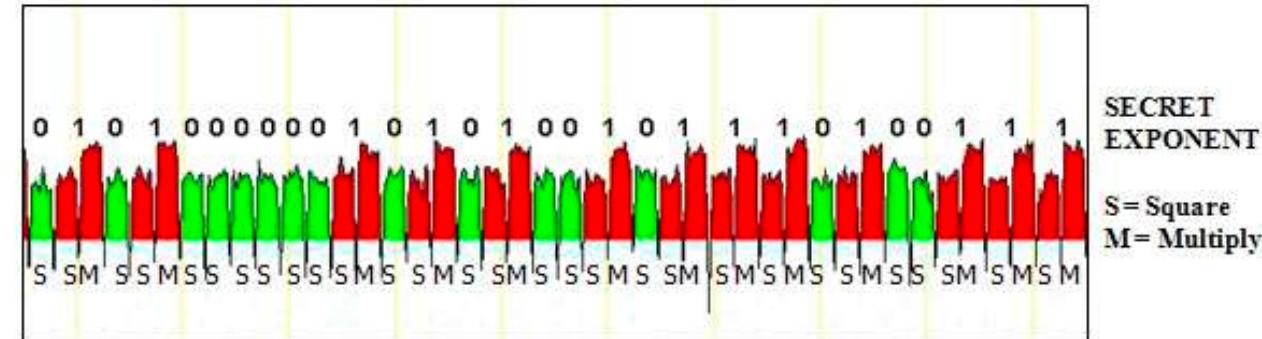
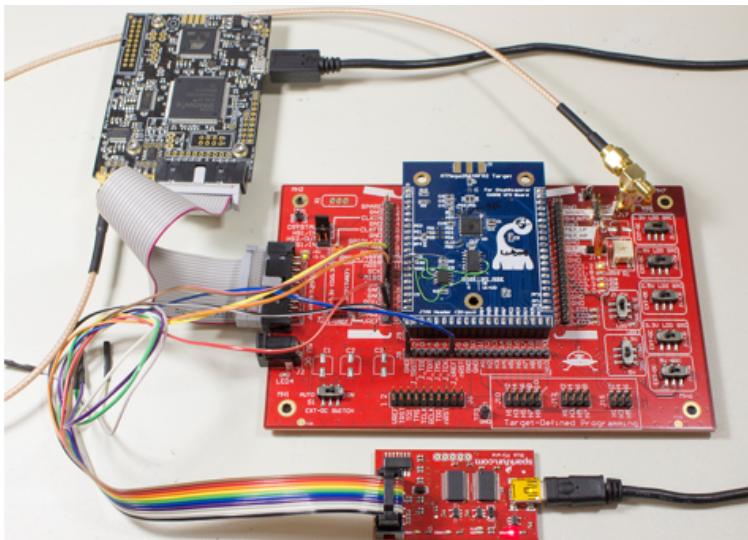
Source: Internet



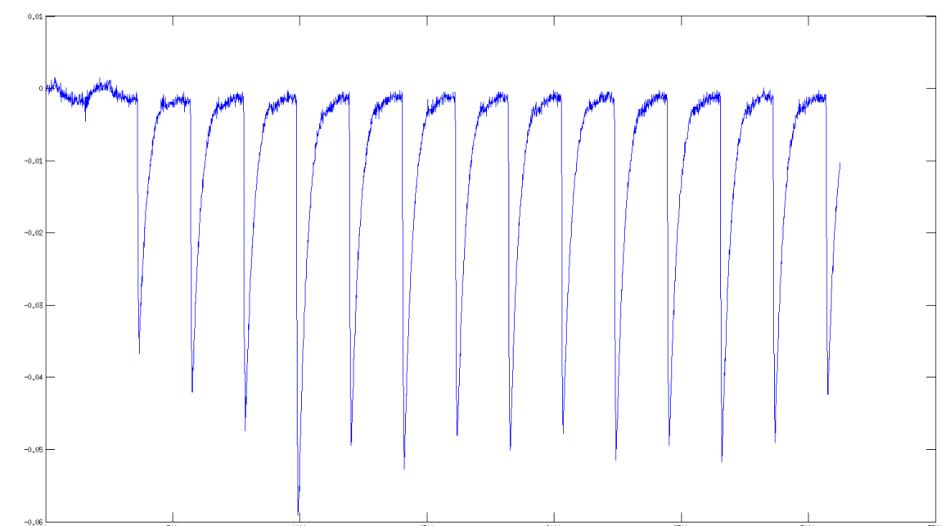
Source: Testbed for Side Channel analysis and security evaluation

# Side-Channel Attacks (SCA)

- The physical channels are correlated with the information being processed
- Fundamental cause: power consumption is correlated with switching of CMOS transistors (0->1, 1->0)
  - Typically it is assumed that power consumption is correlated with the Hamming Weight/ Distance.
- If some internal state is exposed, the secret key can be recovered in seconds.



Source: Internet



Source: Testbed for Side Channel analysis and security evaluation

# Side-Channel Vs. Classical Cryptanalysis

- Cryptanalysis: Purely mathematical
  - Take example of AES
  - Cryptanalysis means, you only have access to plaintext, ciphertext — a lot of them
  - You have to
    - Find the key
    - Or, at least, show that it is distinguishable from uniform randomness

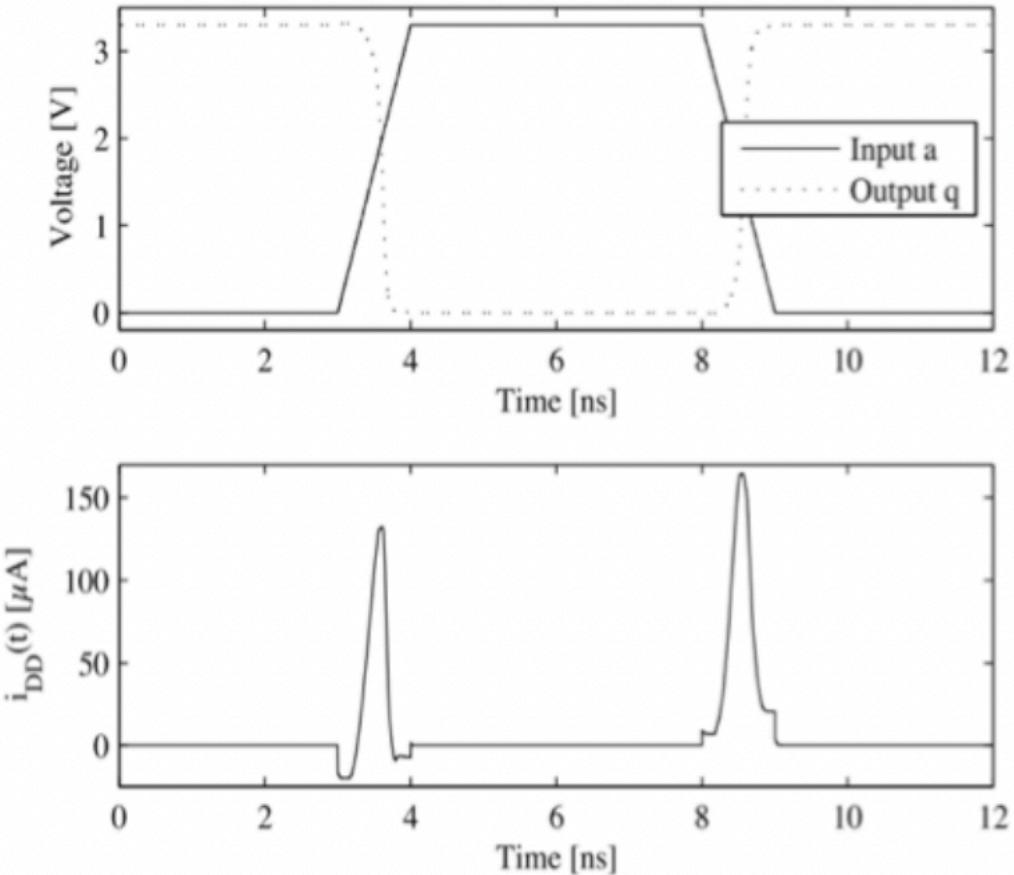
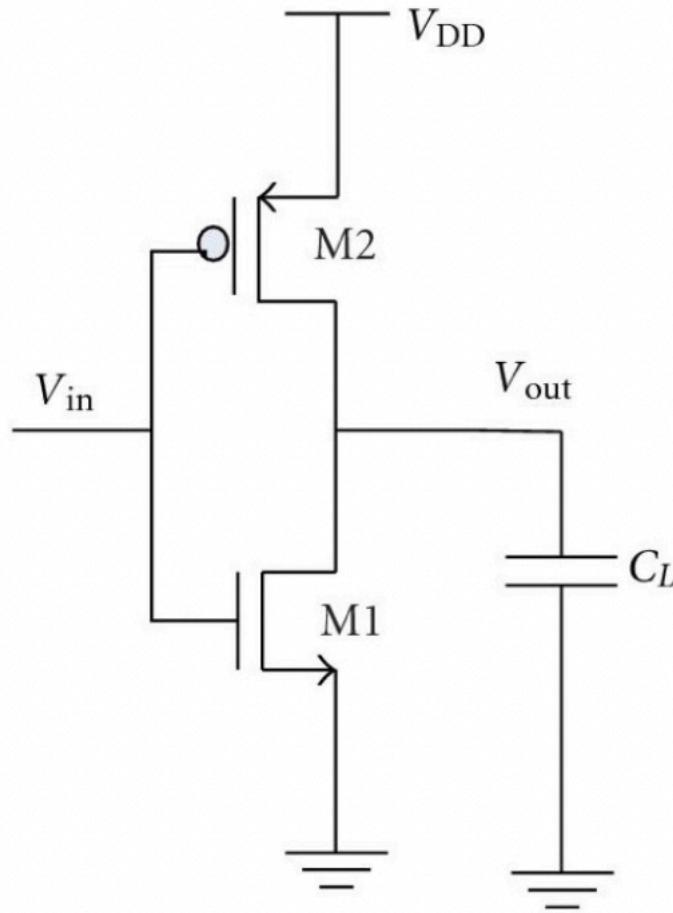
# Side-Channel Vs. Classical Cryptanalysis

- Cryptanalysis: Purely mathematical
  - Take example of RSA/ECC/PQC
  - Cryptanalysis means, you only have access to plaintext, ciphertext — a lot of them
  - You have to
    - Find the key
      - Maybe you need to solve the underlying hard problem in some (mathematical) way.

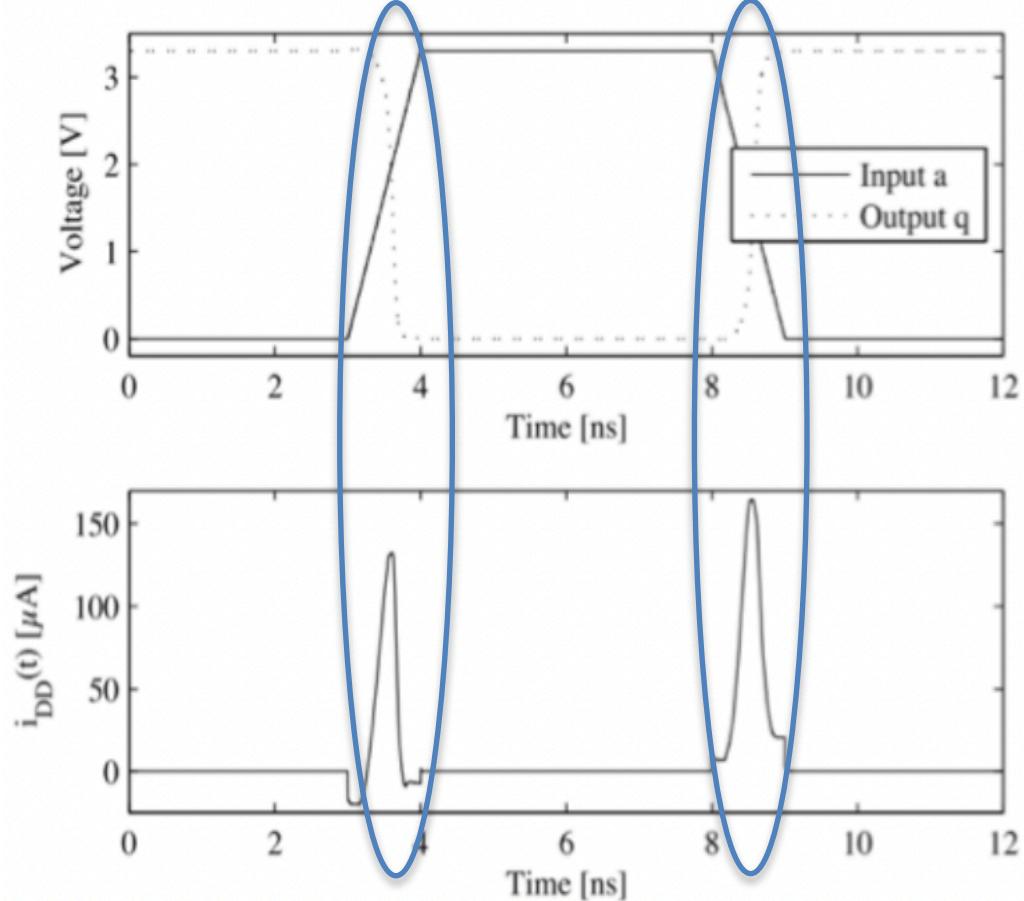
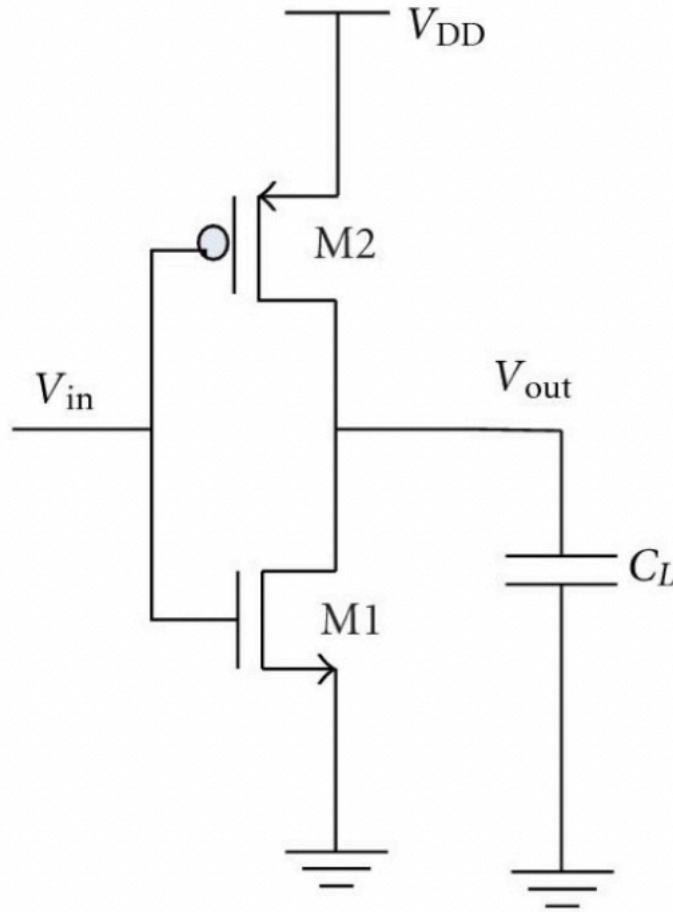
# Side-Channel Vs. Classical Cryptanalysis

- Side-Channel Cryptanalysis: Mathematics + Physics + Statistics
  - The goal is mostly to recover key
    - But also signature forgery, confidentiality breach
  - Ranges beyond crypto...
    - Kernel information extraction
    - Unprivileged access
    - Neural network reverse engineering

# The Root Cause



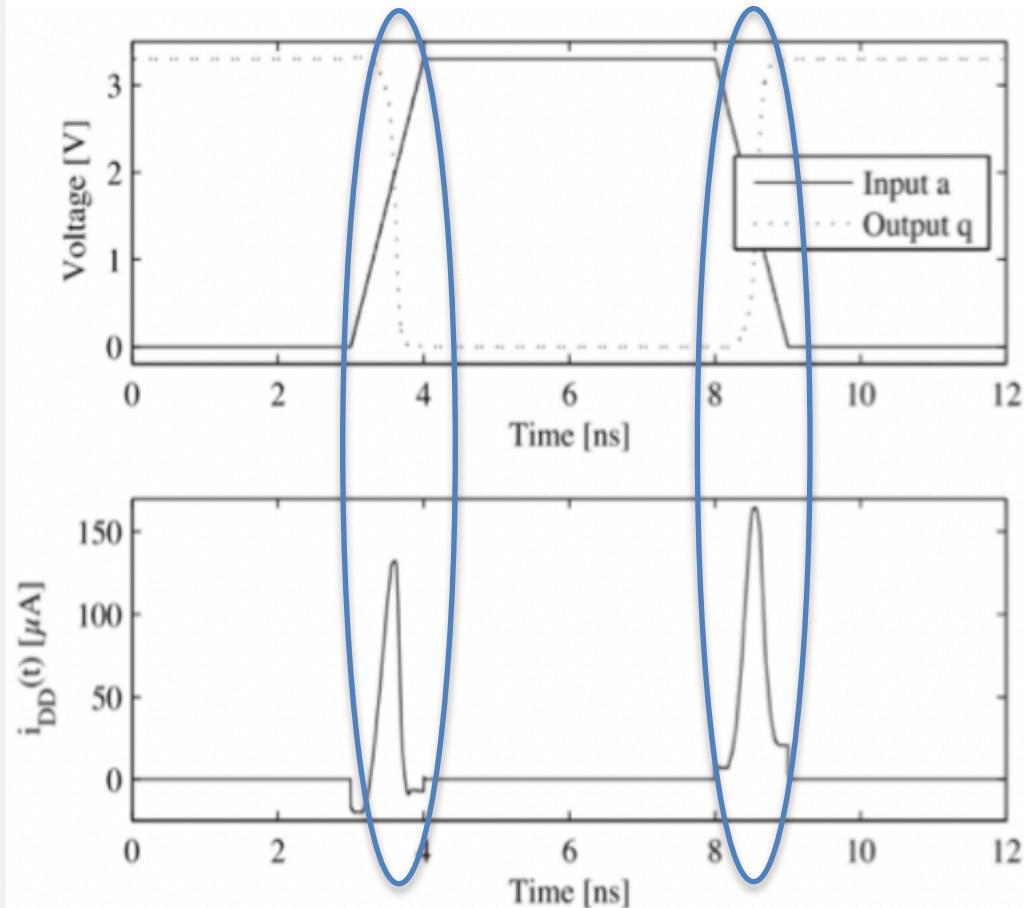
# The Root Cause



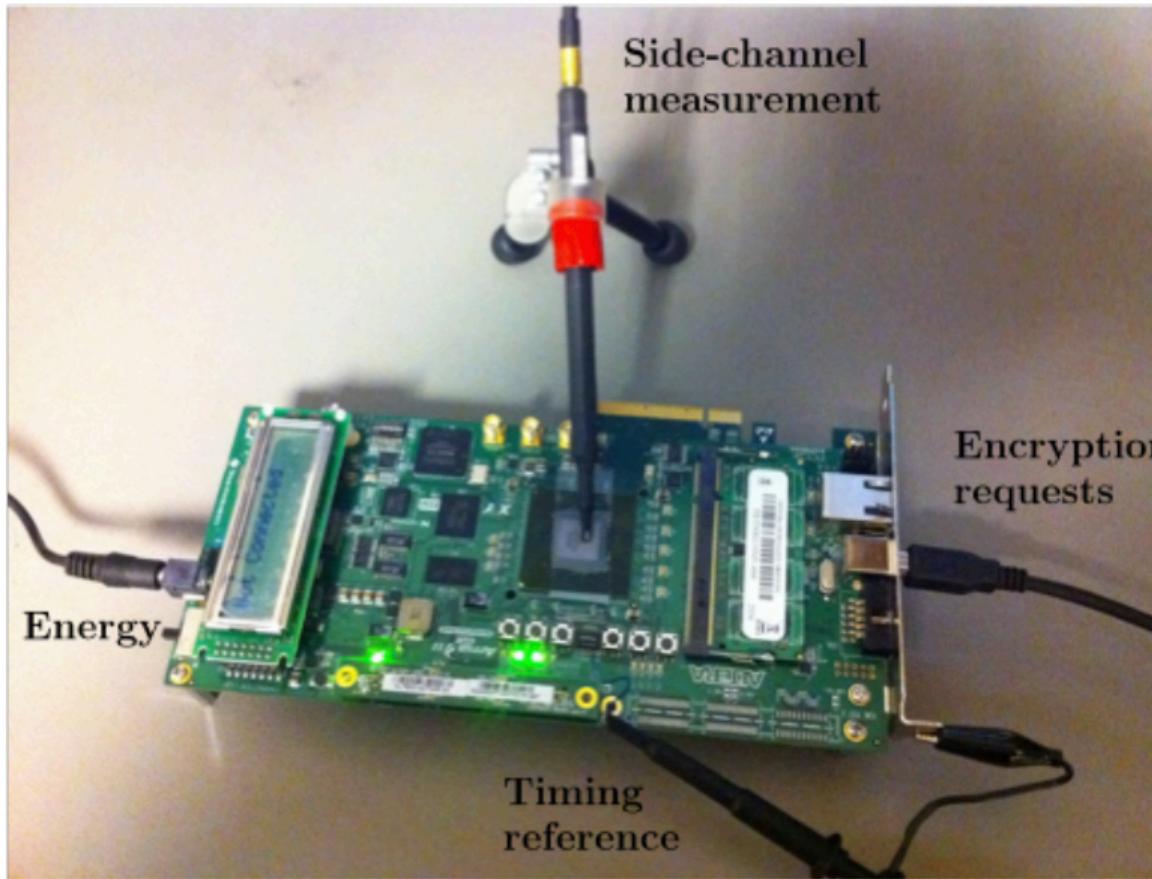
# The Root Cause

## What is exploited?

- The state change of a gate is proportional to the power dissipated.
- Think about a circuit with millions of gates.
- How to measure
  - Power dissipation can be measured by putting a resistor in series with Vdd or Vss and the true source/ground.
  - Roughly, 1 Ohm resistors work well for many microcontrollers, but it is highly target dependent
  - We actually measure current.
    - Differential probes.
  - **The best approach is to use a near-field H-probe an measure EM signal**
    - Less noisy than global power measurement



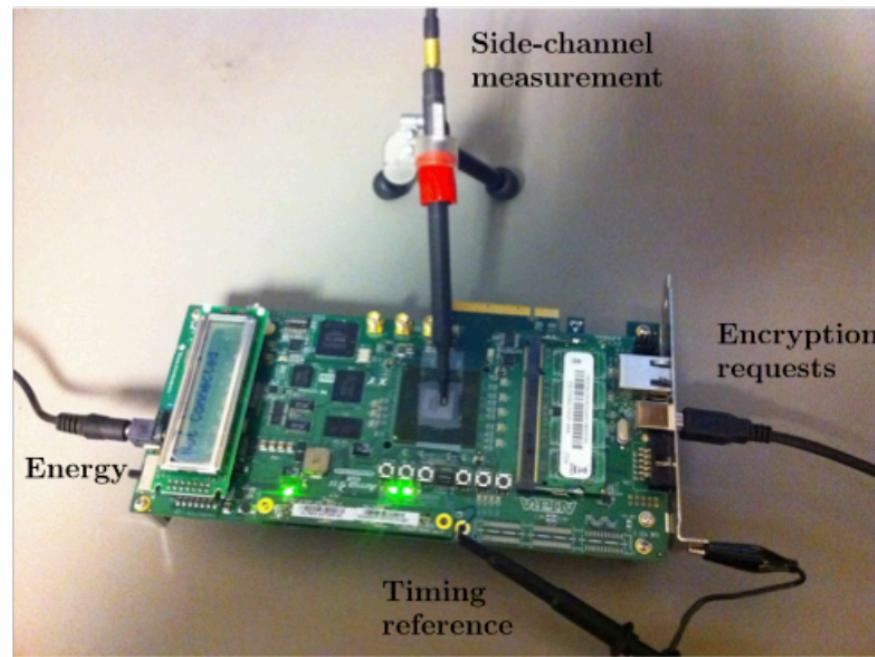
# The Root Cause



# What to “Measure”?

## The Crypto Running on a Microcontroller/ FPGA/ASIC

- End of the day everything is CMOS!!!
- Since power consumption is proportional to the switching activity, so we *can get some idea about the internal computation of the crypto*
  - **The crypto is no more black box**
- In this talk we will be specifically focusing on symmetric key algorithms
  - AES
- What do we mean by attacking AES?
  - Finding out it's secret key



# Looking Inside AES

## State

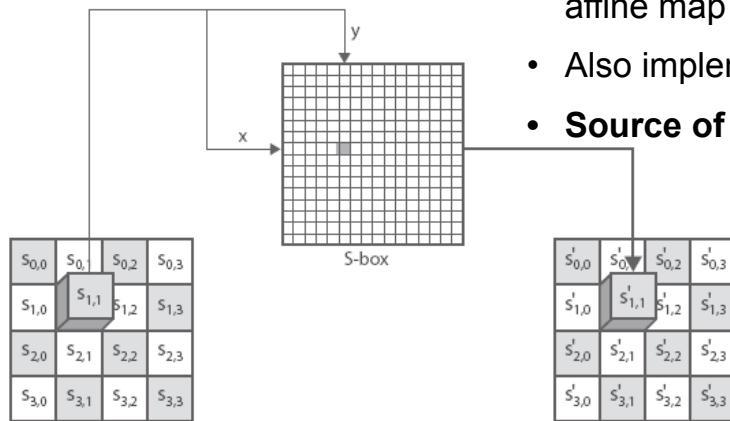
1 byte

|     |     |     |            |
|-----|-----|-----|------------|
| s00 | s01 | s02 | <b>s03</b> |
| s10 | s11 | s12 | s13        |
| s20 | s21 | s22 | s23        |
| s20 | s31 | s32 | s33        |

|     |     |     |             |
|-----|-----|-----|-------------|
| k00 | k01 | k02 | k03 <b></b> |
| k10 | k11 | k12 | k13         |
| k20 | k21 | k22 | k23         |
| k20 | k31 | k32 | k33         |

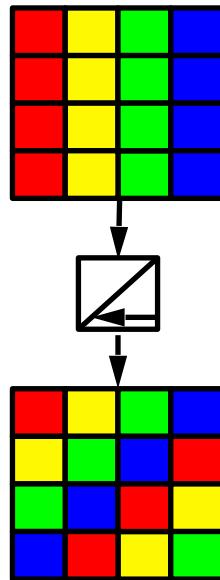
## 1. SubBytes

- Nonlinear Boolean Function
- Finite field inversion followed by affine map
- Also implemented as a table
- **Source of confusion**

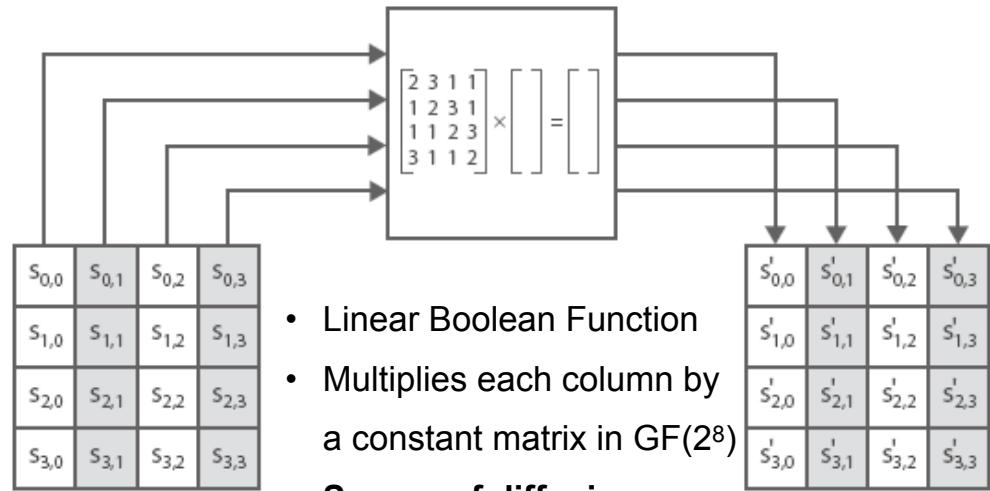


## 2. ShiftRows

- Linear Boolean Function
- Left circular shift of rows
- **Source of diffusion**



## 3. MixColumns



- Linear Boolean Function
- Multiplies each column by a constant matrix in  $\text{GF}(2^8)$
- **Source of diffusion**

## 4. AddRoundKey

|     |     |     |     |
|-----|-----|-----|-----|
| s00 | s01 | s02 | s03 |
| s10 | s11 | s12 | s13 |
| s20 | s21 | s22 | s23 |
| s20 | s31 | s32 | s33 |



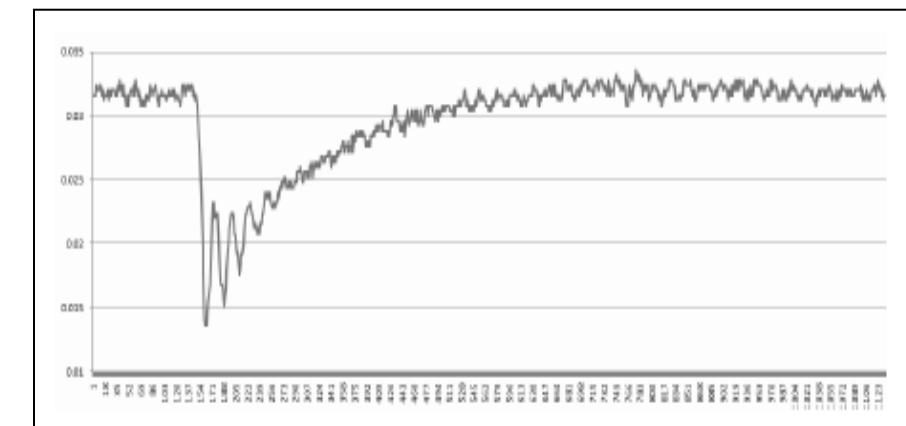
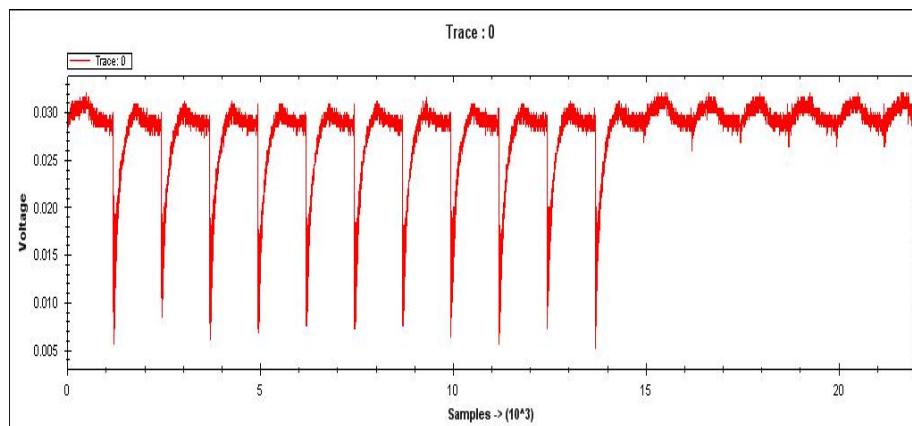
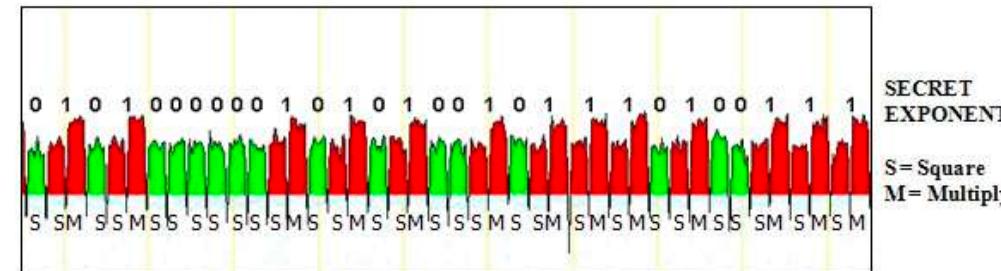
|     |     |     |     |
|-----|-----|-----|-----|
| k00 | k01 | k02 | k03 |
| k10 | k11 | k12 | k13 |
| k20 | k21 | k22 | k23 |
| k20 | k31 | k32 | k33 |

- Linear Boolean Function
- XOR the state with a round key

# Through the Looking Glass

There can be two kinds of power analysis attacks:

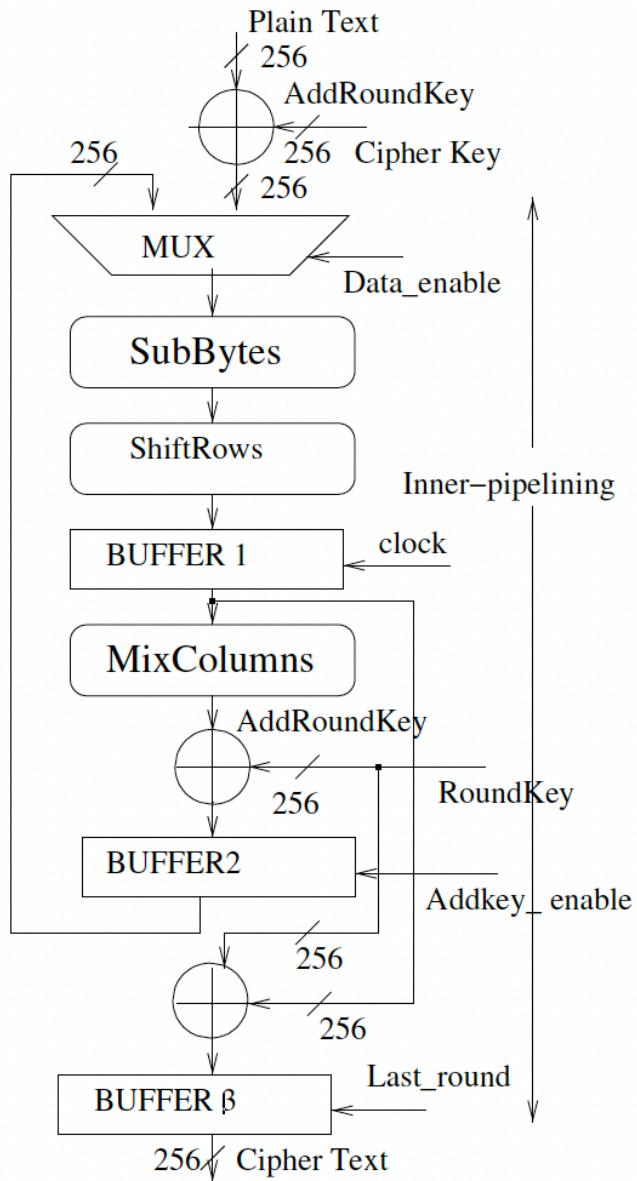
- **Simple power analysis (SPA):** Exploits the operation dependence of power consumption
  - Remember the RSA example from the beginning...
- **Differential Power Analysis (DPA):** Exploits the data dependence of power consumption
  - We will see now for AES
  - **Fact: there is no secret dependent operation in AES, everything is uniform.**



# Differential Power Analysis

- **Power Trace:** A set of power consumptions across a cryptographic process
  - 1 millisecond operation sampled at 5MHz yield a trace with 5000 points.
- **Leakage Model:** Hypothetical model relating the leakage with the internal states of the target algorithm.
  - For AES the internal state is a 128-bit value.
  - **Hamming Weight Model:** The power consumption is proportional to the Hamming weight (count of 1's) of the state.
  - **Hamming Distance Model:** The power consumption is proportional to the Hamming distance between the state in two consecutive clock cycles. — FPGA/ASICs
  - More complex models are possible...
  - Used to simulate leakage and also in some attacks.

# Differential Power Analysis

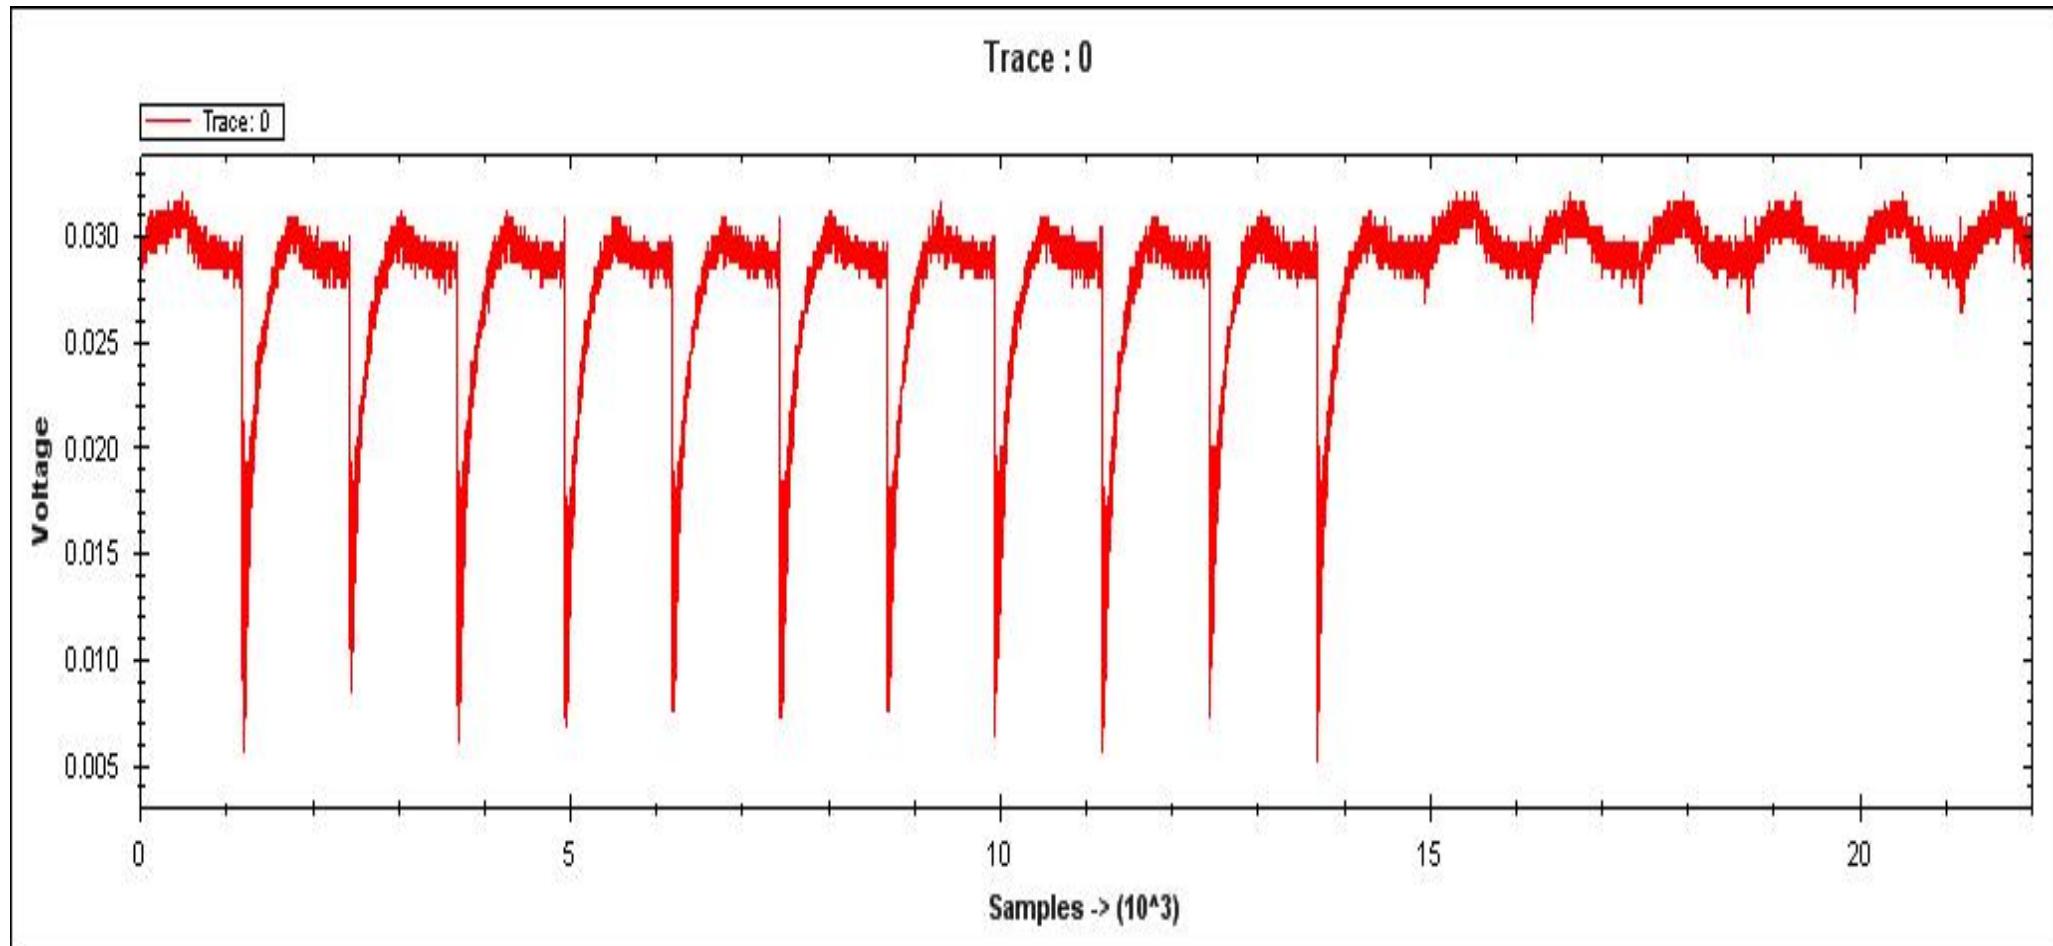


- Common hardware implementation of a block cipher
- We consider the state of the circuit at time instance  $t$  — you can consider it as one time point in the x-axis of the trace.
- Let this state be  $v_t$
- Hamming weight is number of 1's in  $v_t$ .
- Hamming distance is  $\text{HW}(v_t \oplus v_{t-1})$ ..
  - Why?

# Differential Power Analysis

- **Power Trace:** A set of power consumptions across a cryptographic process
  - 1 millisecond operation sampled at 5MHz yield a trace with 5000 points.

- 



# Differential Power Analysis: The Idea

| $s$  | HW( $s$ ) | Target bit (LSB) |
|------|-----------|------------------|
| 0000 | 0         | 0                |
| 0001 | 1         | 1                |
| 0010 | 1         | 0                |
| 0011 | 2         | 1                |
| 0100 | 1         | 0                |
| 0101 | 2         | 1                |
| 0110 | 2         | 0                |
| 0111 | 3         | 1                |
| 1000 | 1         | 0                |
| 1001 | 2         | 1                |
| 1010 | 2         | 0                |
| 1011 | 3         | 1                |
| 1100 | 2         | 0                |
| 1101 | 3         | 1                |
| 1110 | 3         | 0                |
| 1111 | 4         | 1                |

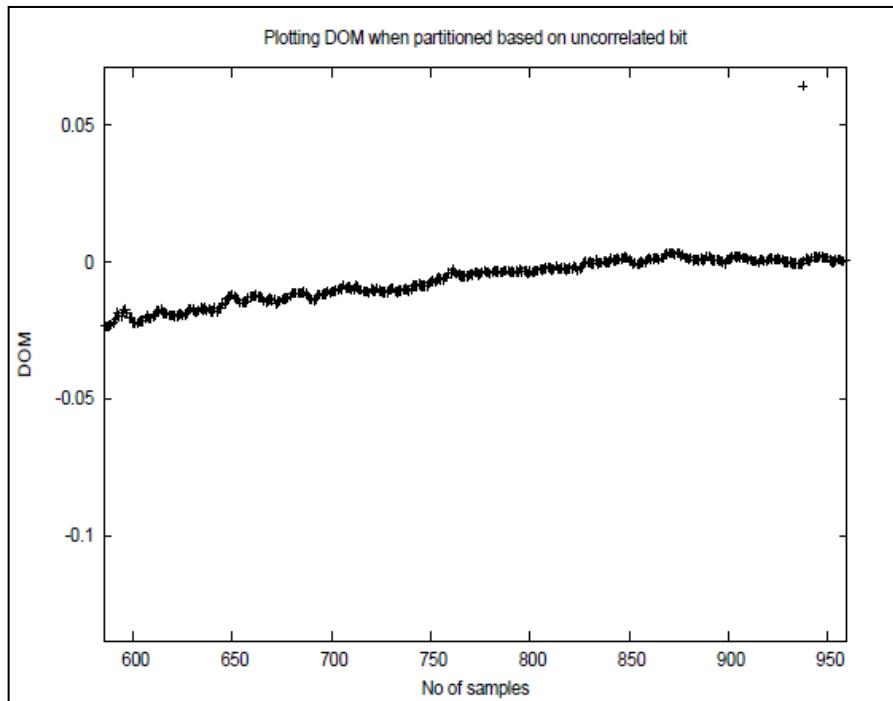
- Assume power leakage follows Hamming Weight.
- Divide the HW( $s$ ) into two bins:
  - 0 bin: when LSB is 0
  - 1 bin: when LSB is 1

# Differential Power Analysis: The Idea

| <i>s</i> | HW( <i>s</i> ) | Target bit (LSB) |
|----------|----------------|------------------|
| 0000     | 0              | 0                |
| 0001     | 1              | 1                |
| 0010     | 1              | 0                |
| 0011     | 2              | 1                |
| 0100     | 1              | 0                |
| 0101     | 2              | 1                |
| 0110     | 2              | 0                |
| 0111     | 3              | 1                |
| 1000     | 1              | 0                |
| 1001     | 2              | 1                |
| 1010     | 2              | 0                |
| 1011     | 3              | 1                |
| 1100     | 2              | 0                |
| 1101     | 3              | 1                |
| 1110     | 3              | 0                |
| 1111     | 4              | 1                |

- Assume power leakage follows Hamming Weight.
- Divide the HW(*s*) into two bins:
  - 0 bin: when LSB is 0
  - 1 bin: when LSB is 1
- Difference-of-Mean (DoM)= $20/8-12/8=1$

# When the Partitioning is Random



- Partitioning done by bits simulated using *rand* function in C.
  - Observe the DoM is close to 0, as expected!
  - **Note:** Instead of computing the difference, we can use some statistical hypothesis test, such as t-test.
  - The hypothesis will be — whether the two trace distributions are the same or different
  - For a random uncorrelated bit, the two distributions are the same.
- 
- Moral of the story: If the bin partitioning is based on a bit from the actual state, there will be significant difference in the mean values of the bins. This is because, the traces are data dependent.

# Attacking AES

## A very very important point

- We haven't seen AES key schedule in detail — but it is has somewhat similar operations as the rounds — has SBoxes, shifts XORs etc.
- **Important:** key schedule is invertible
  - That is, if you recover one round key , you can recover all, and the master key

# Attacking AES

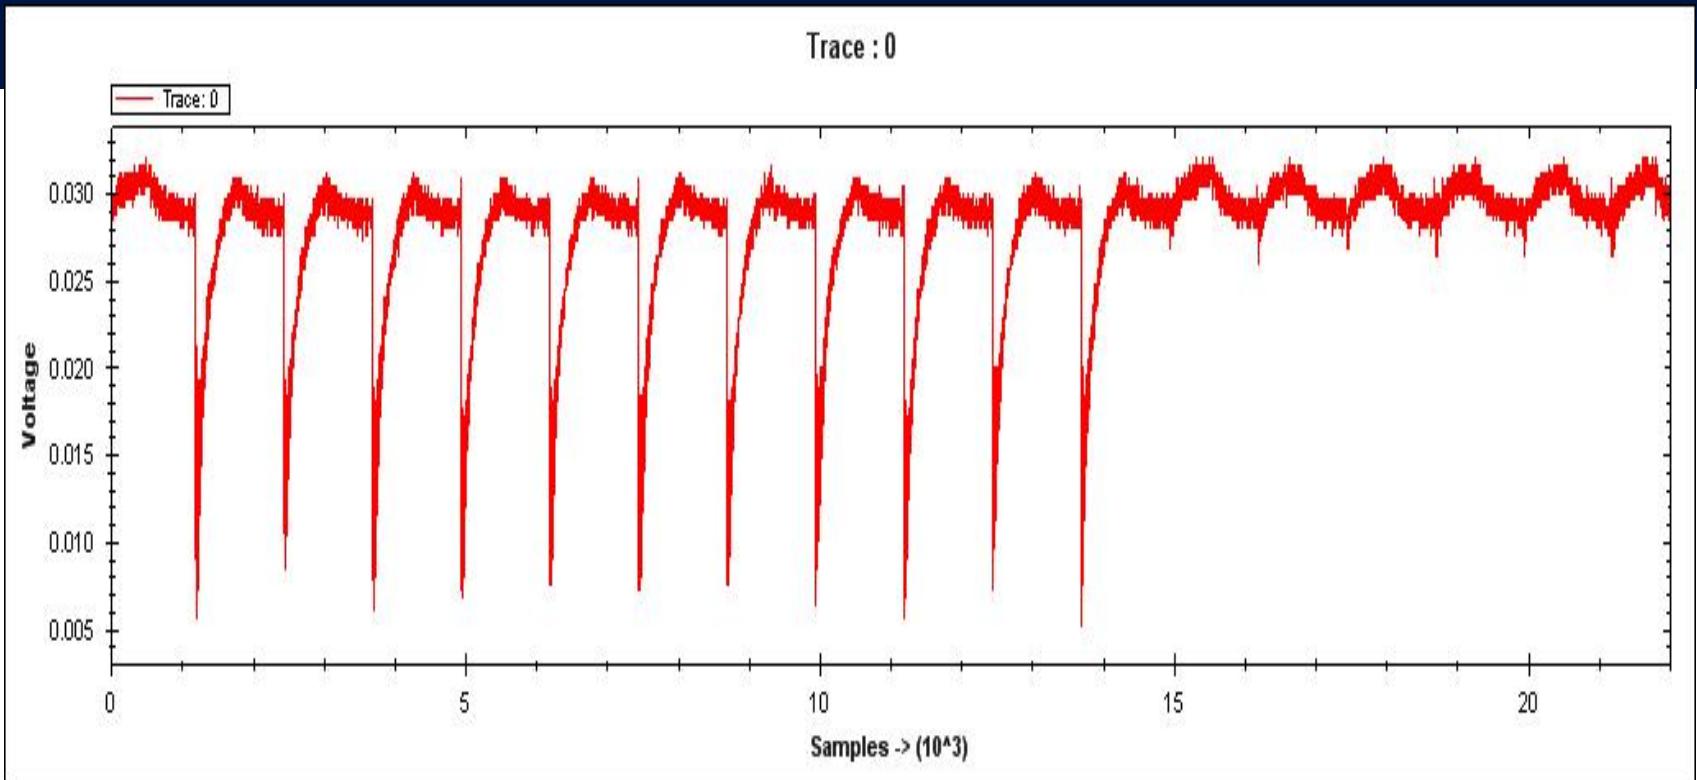
## 1. Measurement:

- Make power consumption measurement of about 1000 AES operations, 100000 data points/trace,
- Save (Ciphertext<sub>i</sub>, trace\_i)

## 2. Attack:

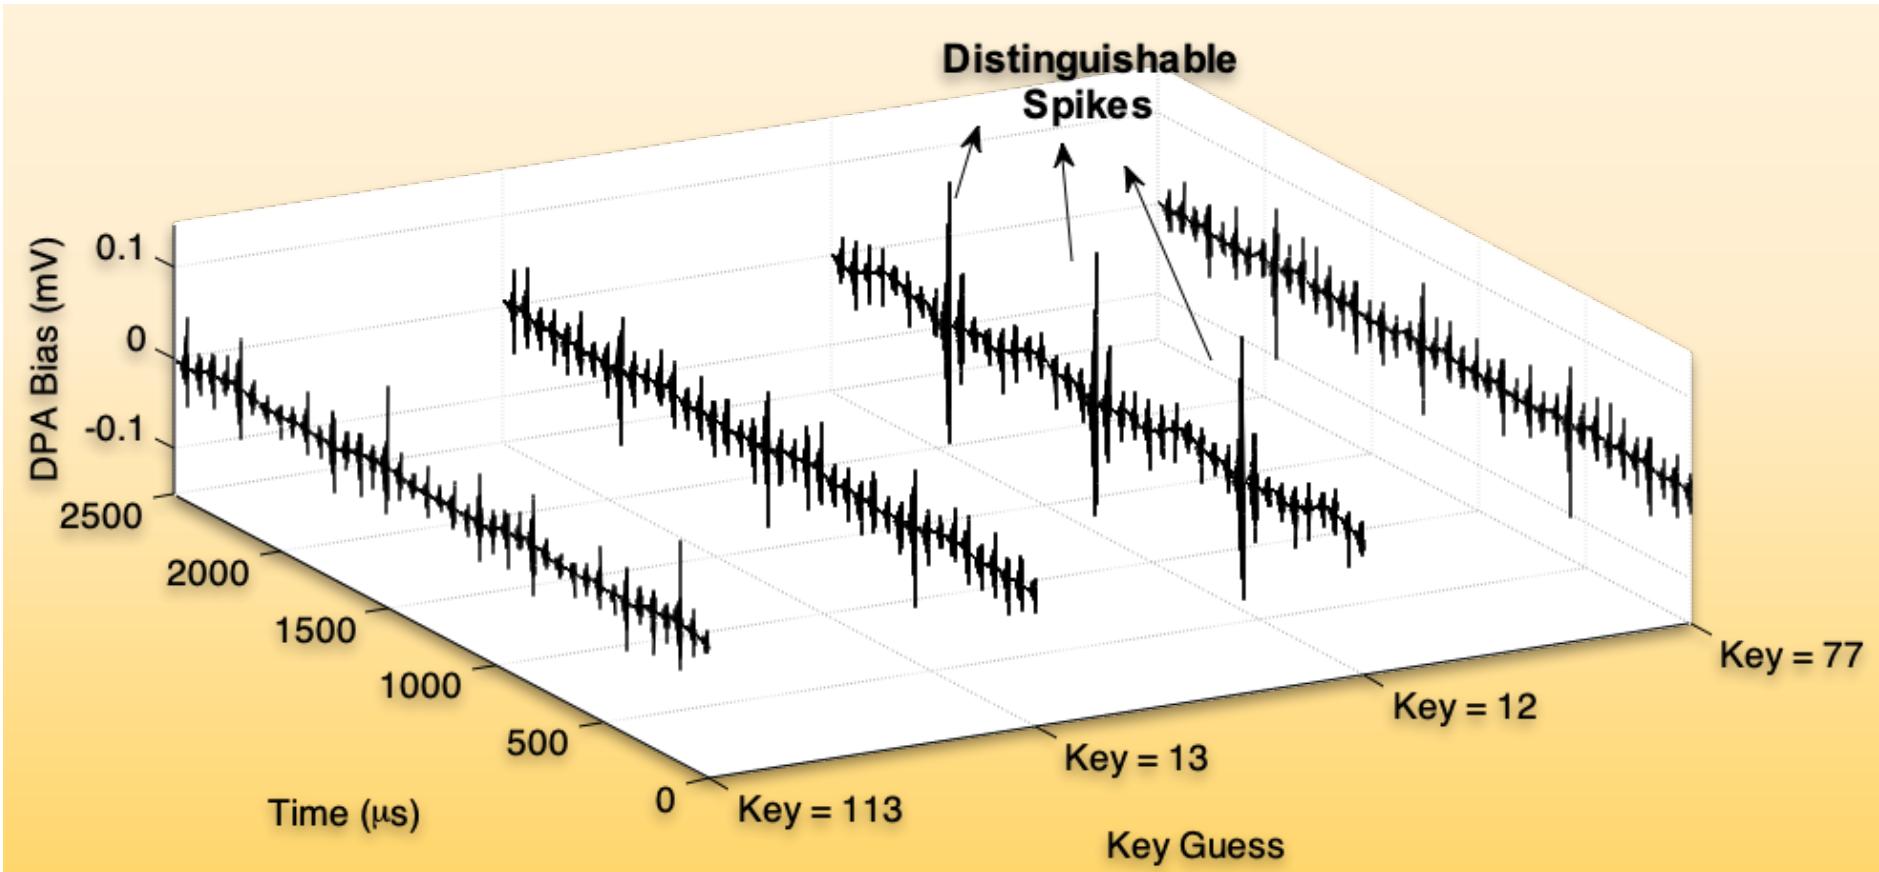
- Target an S-Box in the last round (say the  $j$ -th S-Box)
  - A. Guess a key for an S-box of last round (8 bit key, so total 256 guesses possible)
  - B. Partially decrypt one byte of each ciphertext with the guessed key till the input of the last round S-Box. That is compute:  $S_j = S^{-1}(C_j \oplus k_j^g)$
  - C. Divide the traces into 2 groups based on the LSB of  $S_j$
  - D. Calculate the **average trace** of each group
  - E. Calculate the difference of two **average traces**
  - F. **Correct key guess** → **spikes in the differential curve**
- Repeat A-F for other S-boxes

# Attacking AES



|       |         |         |     |         |
|-------|---------|---------|-----|---------|
| CT[0] | T[0][0] | T[0][1] | ... | T[0][m] |
| CT[1] | T[1][0] | T[1][1] | ... | T[1][m] |
| CT[2] | T[2][0] | T[2][1] | ... | T[2][m] |
| ...   | ...     | ...     | ... | ...     |
| CT[n] | T[n][0] | T[n][1] | ... | T[n][m] |

# Attacking AES



SBOX – 11

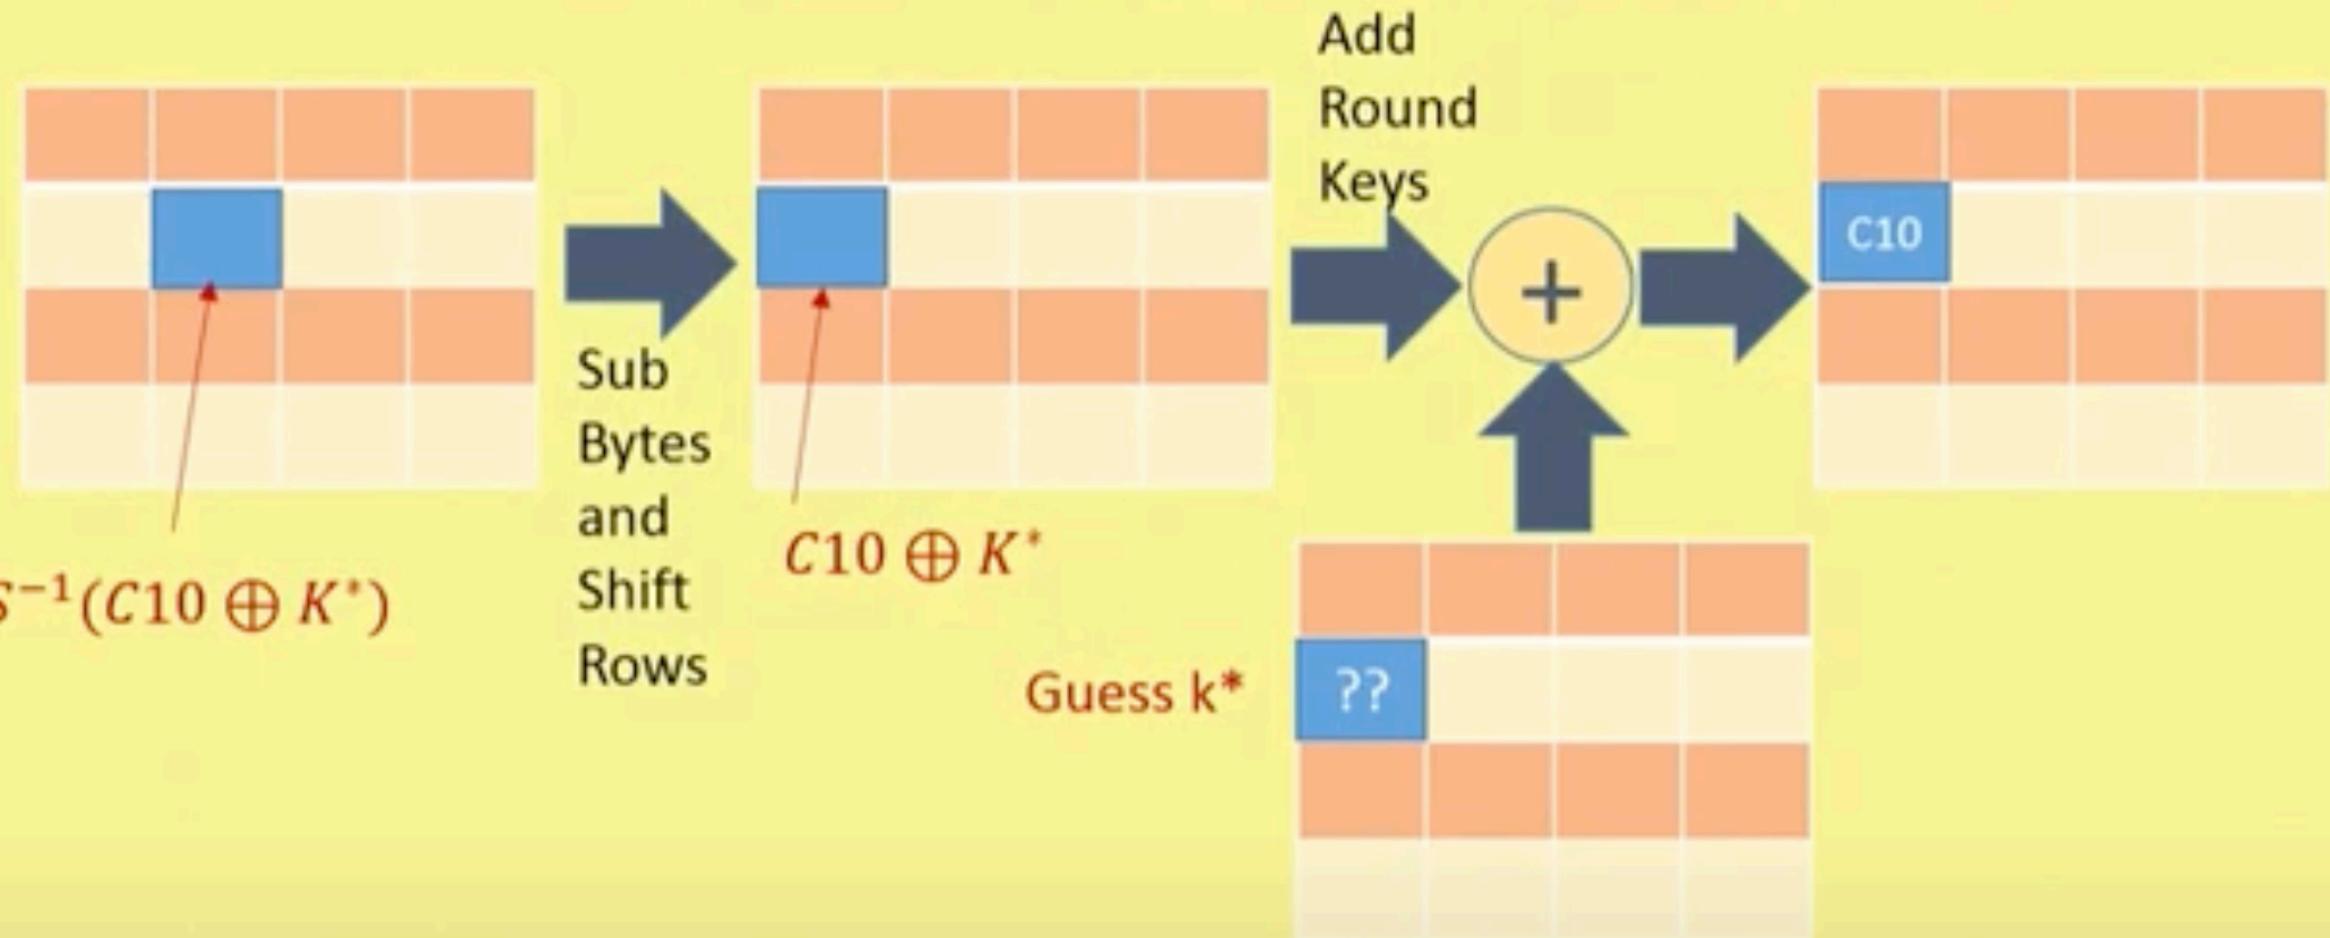
BIT – 8 TRACE COUNT = 15,000, FPGA implementation

# Attacking AES

- **DPA selection function:**  $D(C, b, k^g)$  is defined as computing the value of the
  - $b^{\text{th}}$  output bit, depending upon
    - C: Ciphertext
    - $k^g$  is the guessed key for the S-Box
- In the attack,  $D = S^{-1}(C_j \oplus k_j^g)$
- If  $k^g$  is a wrong guess then  $b$  is correctly evaluated only for half of the ciphertexts (randomly).
  - **Thus for large number of points, the difference between average traces is close to 0**
  - **In other words, distribution of both the bins will be the same**
- But if  $k^g$  is a correct guess, then  $b$  is correctly evaluated for all the ciphertexts.

**Principle:** If  $K_s$  is wrongly guessed, D behaves like a random guess. Thus for a large number of sample points,  $\Delta[1..k]$  tends to zero. But if its correct, the differential will be non-zero and show spikes when D is correlated with the value being processed.

# Attacking AES



$$D(C10, b = 0, K10) = S^{-1}(C10 \oplus K^*)|_{b=0}$$

# Attacking AES

- **Differential Trace:** It is a  $m$  sample trace denoted as  $\Delta_D$ , where,

$$\Delta_D[j] = \frac{\sum_{i=0}^{n-1} D(C_i, b, K_g) T[i][j]}{\sum_{i=0}^{n-1} D(C_i, b, K_g)} - \frac{\sum_{i=0}^{n-1} (1 - D(C_i, b, K_g)) T[i][j]}{\sum_{i=0}^{n-1} (1 - D(C_i, b, K_g))}$$

- **Note:**  $C_i$  is a particular byte of the  $i$ -th ciphertext.

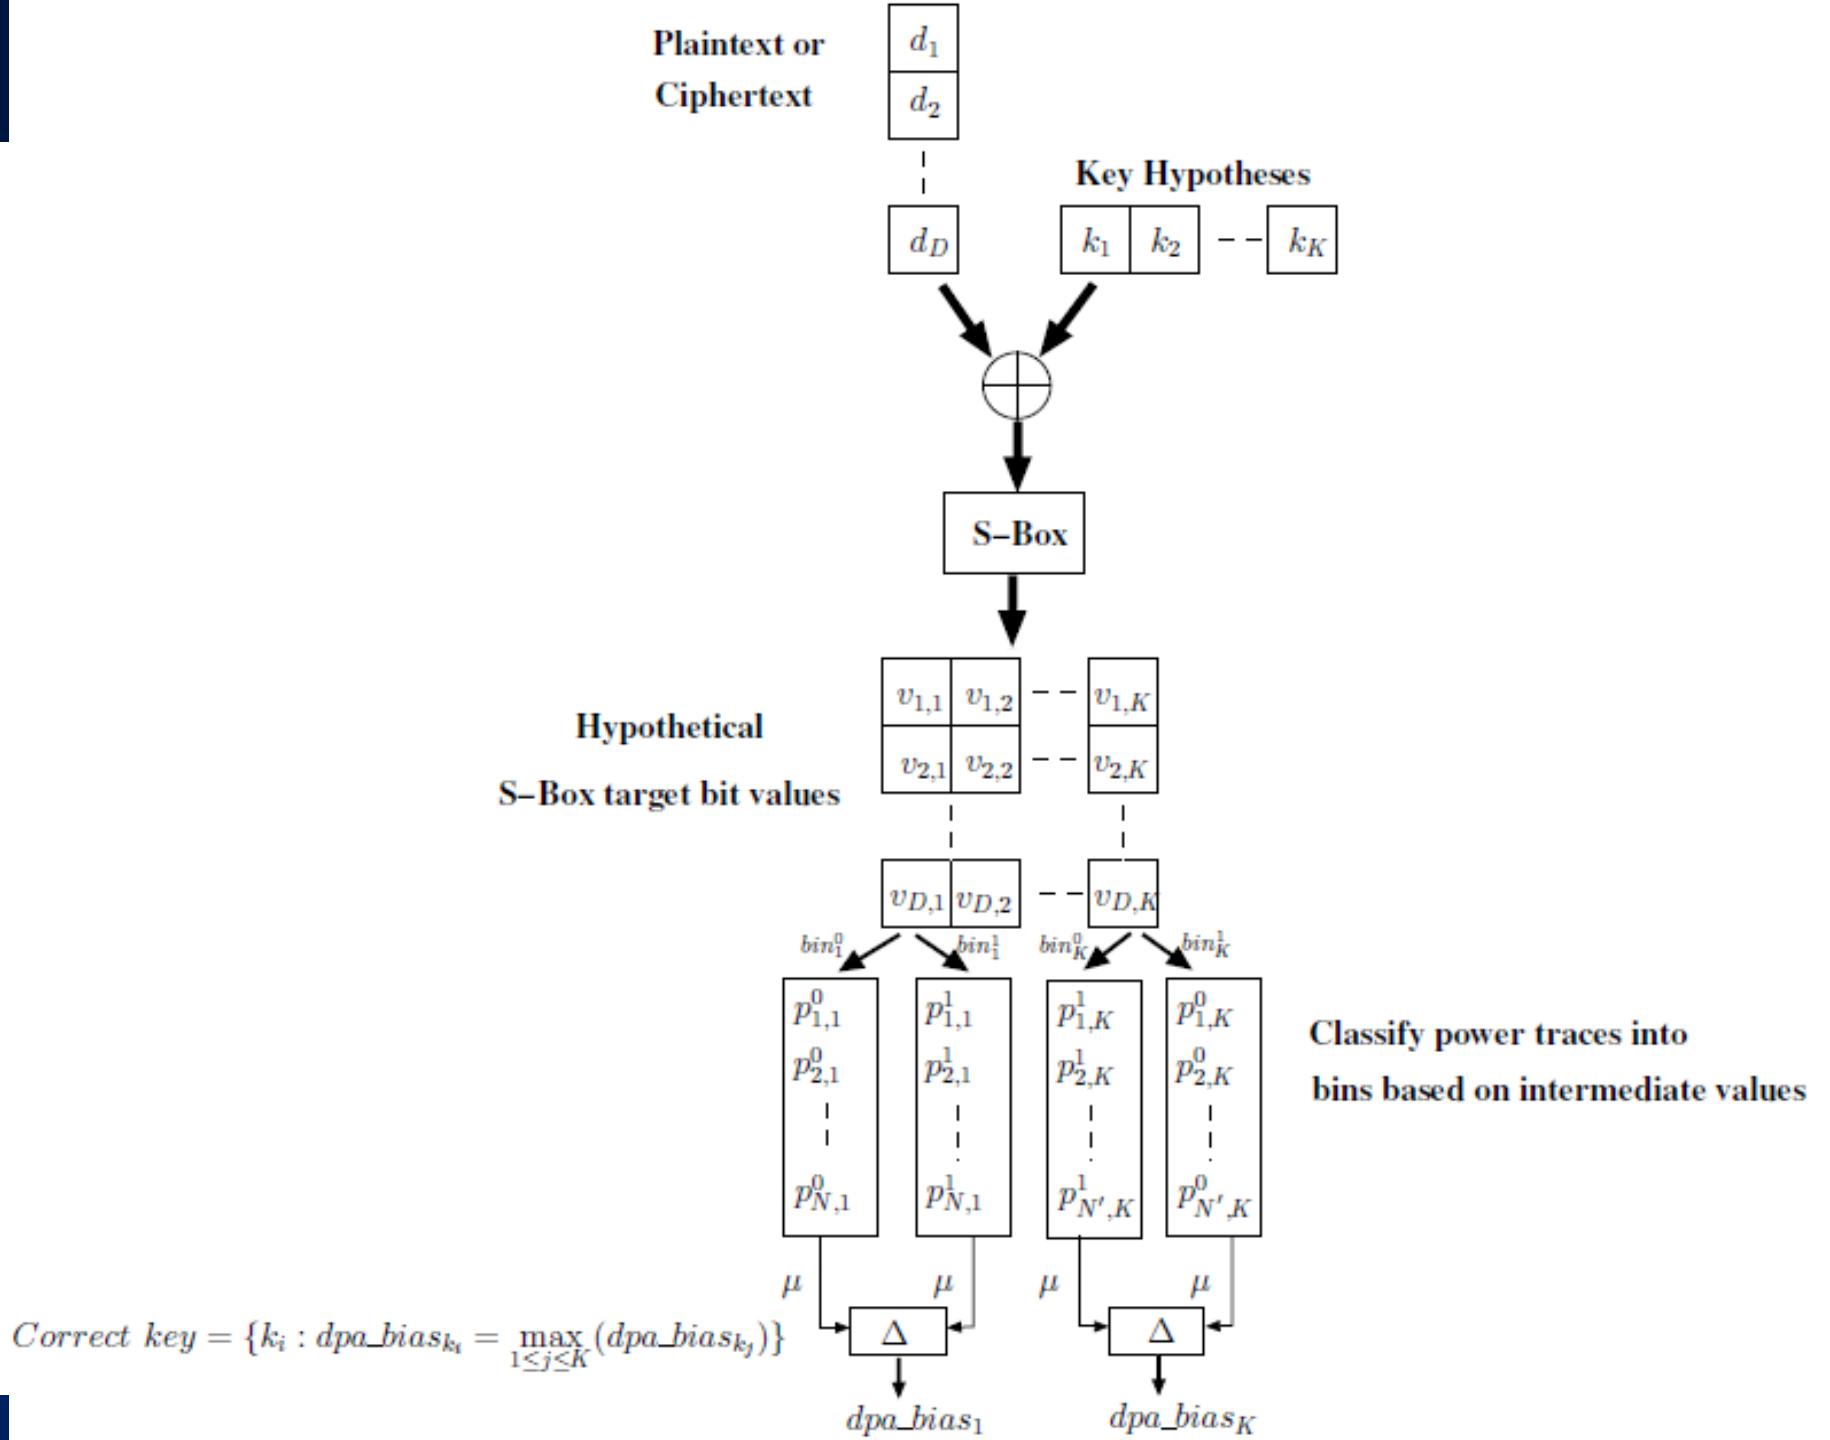
# Attacking AES

- **Why does the attack work?**

- It's not only the data dependency. But it also depends on the mathematics of AES
- **DPA selection function:**  $D(C, b, k^g)$  is defined as computing the value of the
  - $b^{\text{th}}$  output bit, depending upon
    - C: Ciphertext
    - $k^g$  is the guessed key for the S-Box
- In the attack,  $D = S^{-1}(C_j \oplus k_j^g)$
- If  $k^g$  is a wrong guess then  $b$  is correctly evaluated only for half of the ciphertexts (randomly).
  - **Thus for large number of points, the difference between average traces is close to 0**
  - **In other words, distribution of both the bins will be the same**
- But if  $k^g$  is a correct guess, then  $b$  is correctly evaluated for all the ciphertexts.
- **Note:** The non-linearity of the S-Boxes play an important role here.

**Principle:** If  $K_s$  is wrongly guessed, D behaves like a random guess. Thus for a large number of sample points,  $\Delta[1..k]$  tends to zero. But if its correct, the differential will be non-zero and show spikes when D is correlated with the value being processed.

# Attacking AES



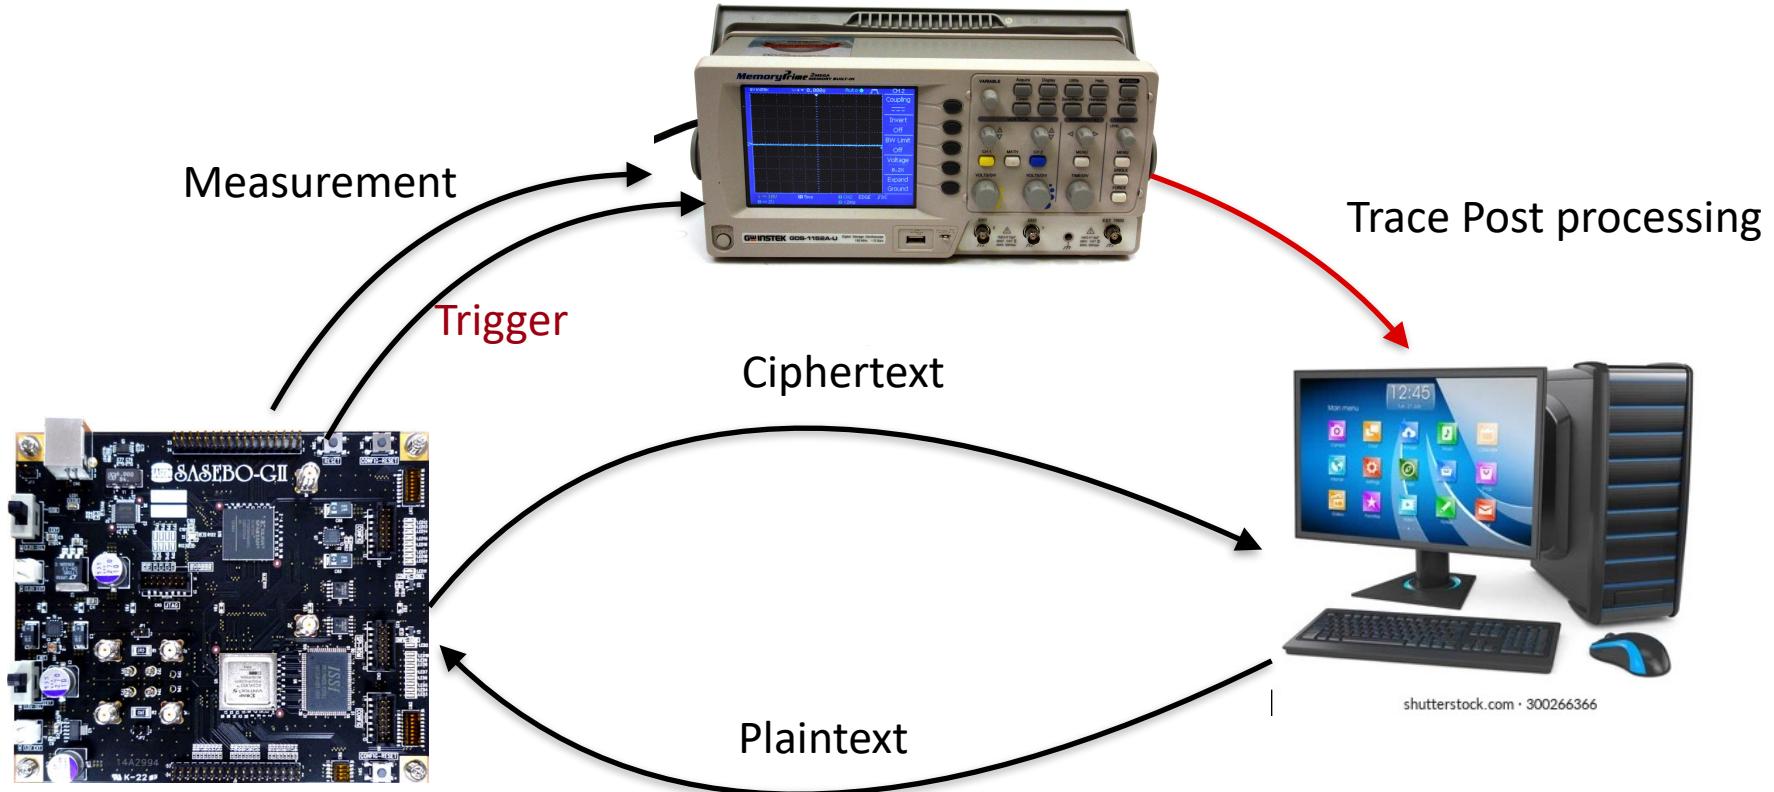
# Attacking AES: Attack Complexity

- **What is the attack complexity**
  - Say you are given n number of traces
  - How much further computation you need to perform?

# Attacking AES: Attack Complexity

- **Observe:** We are targeting the last round of AES and attacking each S-Box separately.
  - Divide and Conquer
  - Attack complexity: always  $2^8$
  - Trace complexity: depends
- **Observe:** The attack don't need plaintext knowledge

# Leakage Models



shutterstock.com · 300266366

# Differential Power Analysis

- **Power Trace:** A set of power consumptions across a cryptographic process
  - 1 millisecond operation sampled at 5MHz yield a trace with 5000 points.
- **Measurement:**
  - Usually requires to adhere a signal bandwidth and frequency.
    - We do quite high frequency sampling — e.g. for a 65 Mhz microcontroller or 150 Mhz FPGA implementation, we do sampling at 5 GS/sec.
  - Number of samples is an important point. More samples results in better attack
  - **Trigger:** An important constraint
    - Without reliable trigger, the traces will be **misaligned**.— attack will be difficult
    - Most of the practical implementations does not provide a very reliable trigger
      - So, some realignment of traces are needed
        - Today, DL algorithms are great for this purpose — does automatic realignment.

# Attacking AES

## 1. Measurement:

- Make power consumption measurement of about 1000 AES operations, 100000 data points/trace,
- Save (Ciphertext<sub>i</sub>, trace\_i)

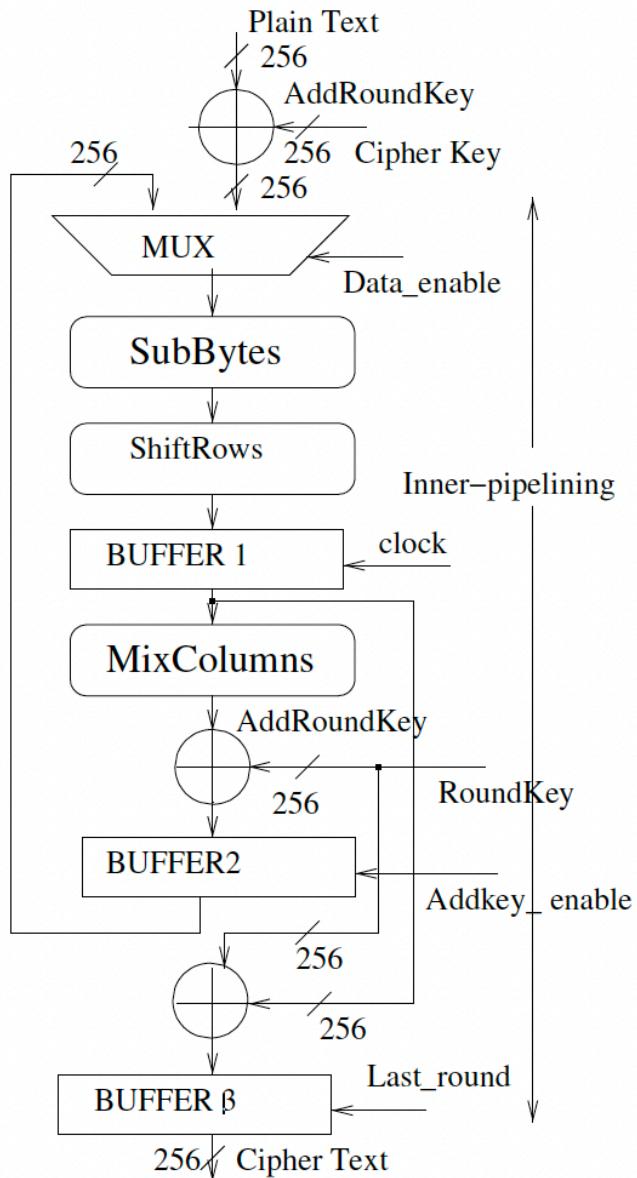
## 2. Attack:

- Target an S-Box in the last round (say the  $j$ -th S-Box)
  - A. Guess a key for an S-box of last round (8 bit key, so total 256 guesses possible)
  - B. Partially decrypt one byte of each ciphertext with the guessed key till the input of the last round S-Box. That is compute:  $S_j = S^{-1}(C_j \oplus k_j^g)$
  - C. Divide the traces into 2 groups based on the LSB of  $S_j$
  - D. Calculate the **average trace** of each group
  - E. Calculate the difference of two **average traces**
  - F. **Correct key guess** → **spikes in the differential curve**
- Repeat A-F for other S-boxes

# Differential Power Analysis

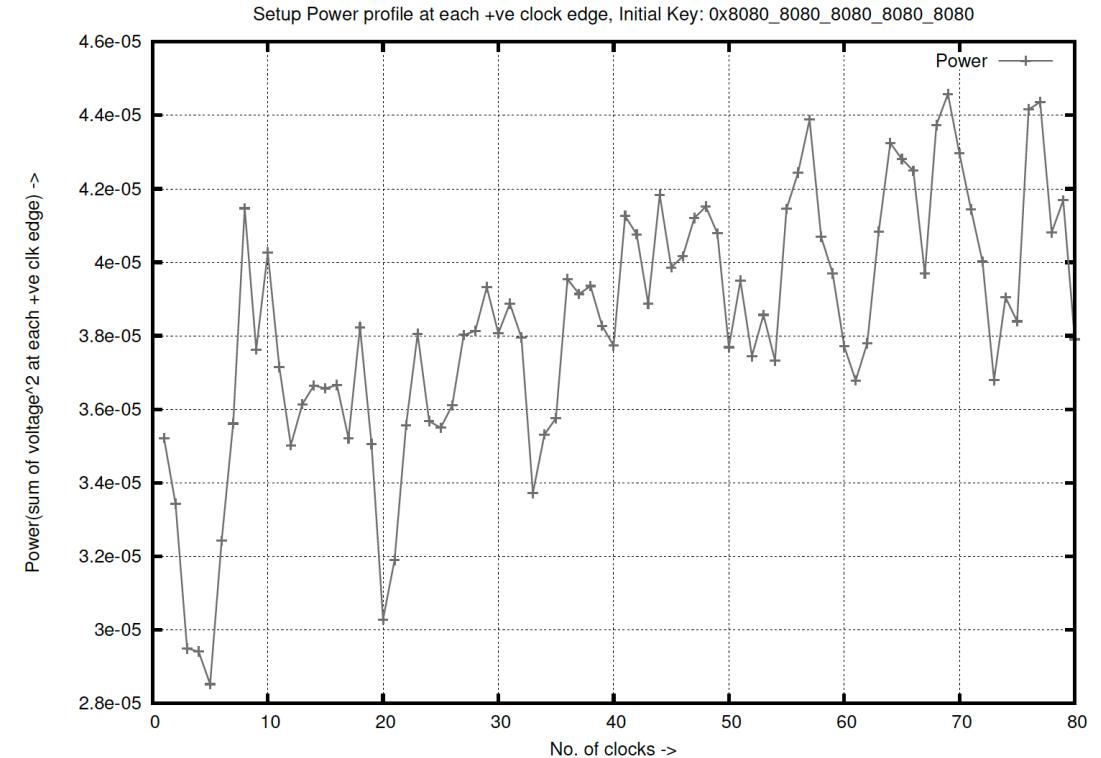
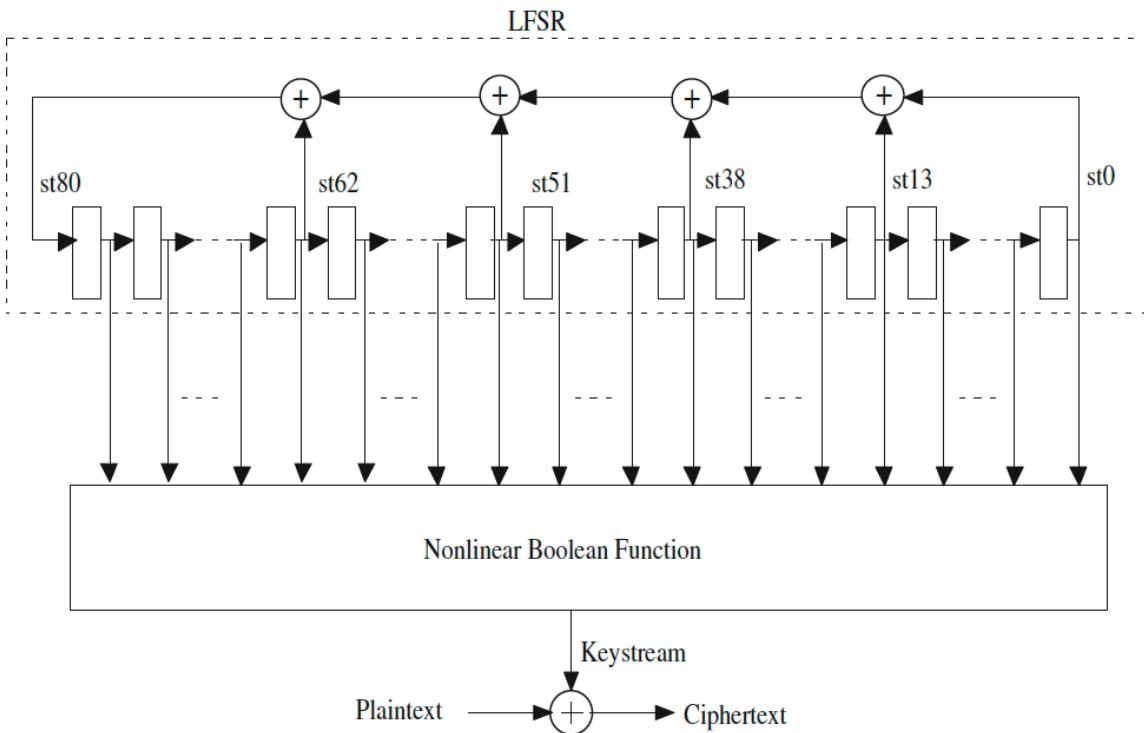
- **Power Trace:** A set of power consumptions across a cryptographic process
  - 1 millisecond operation sampled at 5MHz yield a trace with 5000 points.
- **Leakage Model:** Hypothetical model relating the leakage with the internal states of the target algorithm.
  - For AES the internal state is a 128-bit value.
  - **Hamming Weight Model:** The power consumption is proportional to the Hamming weight (count of 1's) of the state.
  - **Hamming Distance Model:** The power consumption is proportional to the Hamming distance between the state in two consecutive clock cycles. — FPGA/ASICs
  - More complex models are possible...
  - Used to simulate leakage and also in some attacks.

# Differential Power Analysis



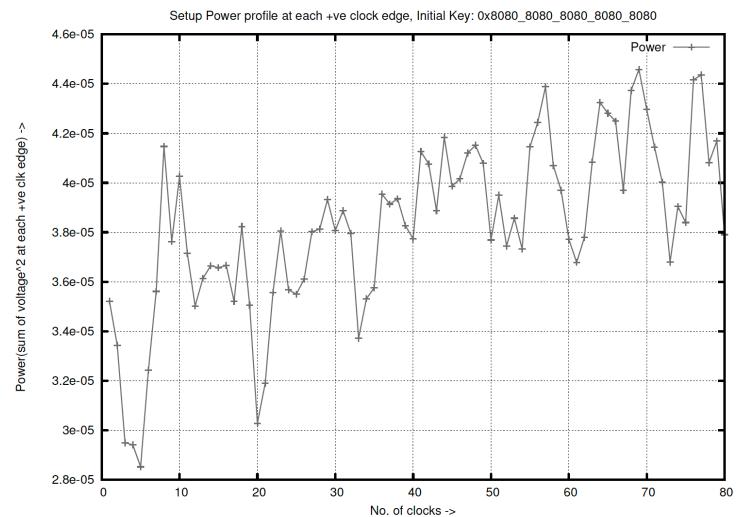
- Common hardware implementation of a block cipher
- We consider the state of the circuit at time instance  $t$  — you can consider it as one time point in the x-axis of the trace.
- Let this state be  $v_t$
- Hamming weight is number of 1's in  $v_t$ .
- Hamming distance is  $\text{HW}(v_t \oplus v_{t-1})$ .
  - Why?
  - After all power consumption is a result of transition  $v_{t-1} \rightarrow v_t$
  - So, HW is relatively inaccurate, but works well for software
    - Registers often start from a specific initialisation value.

# Differential Power Analysis

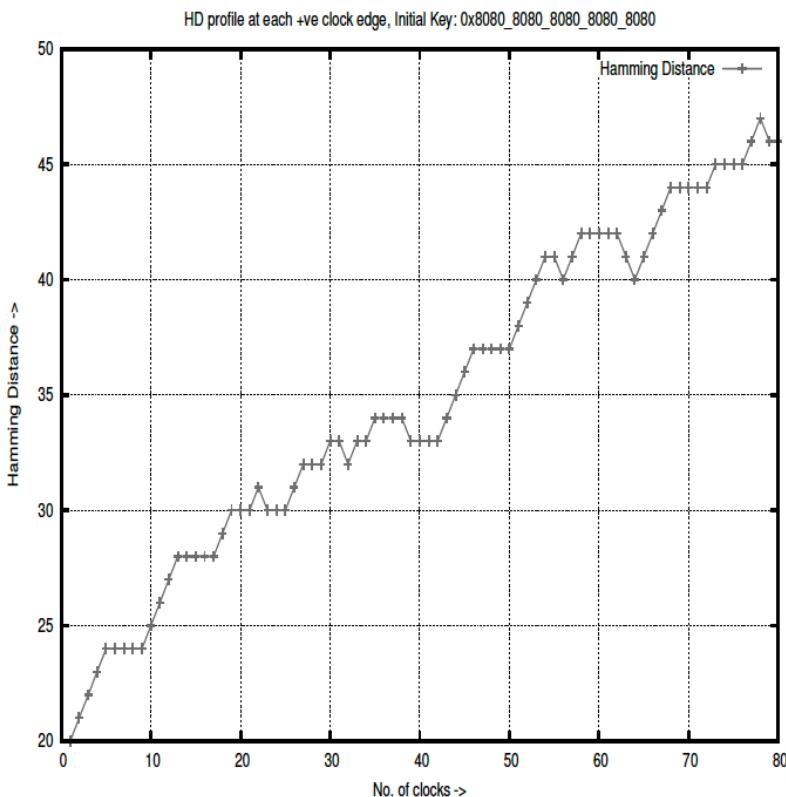


Actual Trace from an FPGA

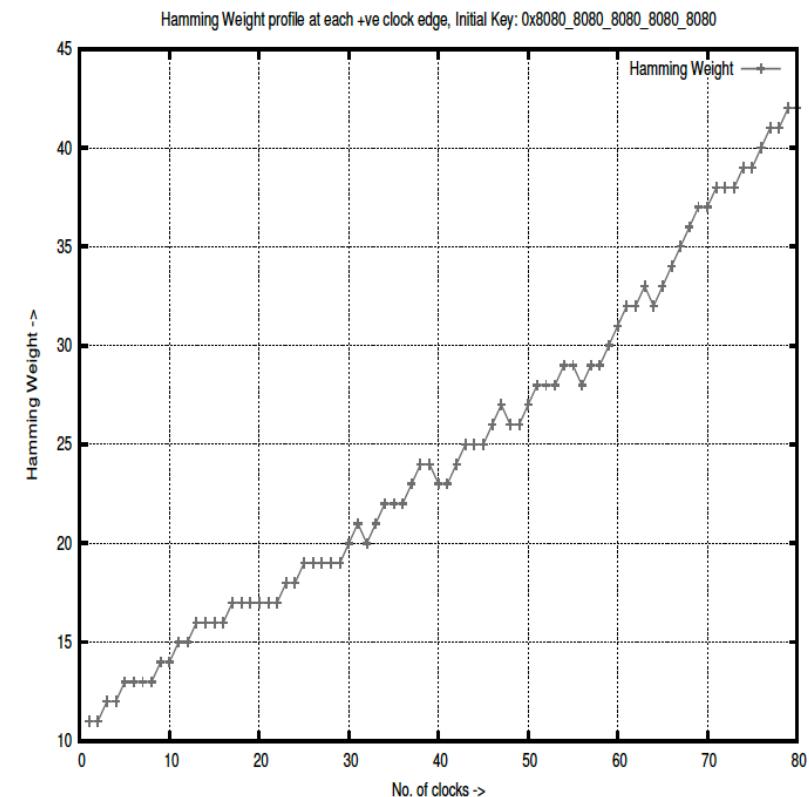
# Differential Power Analysis



Actual Trace from :



Actual Trace from an FPGA



# Why do Gates Leak?

| a   | b   | y   | Energy    |
|-----|-----|-----|-----------|
| 0→0 | 0→0 | 0→0 | $E_{0→0}$ |
| 0→0 | 0→1 | 0→0 | $E_{0→0}$ |
| 0→0 | 1→0 | 0→0 | $E_{0→0}$ |
| 0→0 | 1→1 | 0→0 | $E_{0→0}$ |
| 0→1 | 0→0 | 0→0 | $E_{0→0}$ |
| 0→1 | 0→1 | 0→1 | $E_{0→1}$ |
| 0→1 | 1→0 | 0→0 | $E_{0→0}$ |
| 0→1 | 1→1 | 0→1 | $E_{0→1}$ |
| 1→0 | 0→0 | 0→0 | $E_{0→0}$ |
| 1→0 | 0→1 | 0→0 | $E_{0→0}$ |
| 1→0 | 1→0 | 1→0 | $E_{1→0}$ |
| 1→0 | 1→1 | 1→0 | $E_{1→0}$ |
| 1→1 | 0→0 | 0→0 | $E_{0→0}$ |
| 1→1 | 0→1 | 0→1 | $E_{0→1}$ |
| 1→1 | 1→0 | 1→0 | $E_{1→0}$ |
| 1→1 | 1→1 | 1→1 | $E_{1→1}$ |

- Consider an AND gate  $y = ab$
- 4 different energy levels due to transition of the gate  $E_{0→0}, E_{0→1}, E_{1→0}, E_{1→1}$
- We estimate  $E(q = 0)$ , and  $E(q = 1)$  which are the average energy levels when  $q = 0$ , and  $q = 1$ , respectively.
- $E(q = 0) = \frac{9E_{0→0} + 3E_{1→0}}{12}$
- $E(q = 1) = \frac{3E_{0→1} + E_{1→1}}{4}$
- So  $E(q = 0) \neq E(q = 1)$

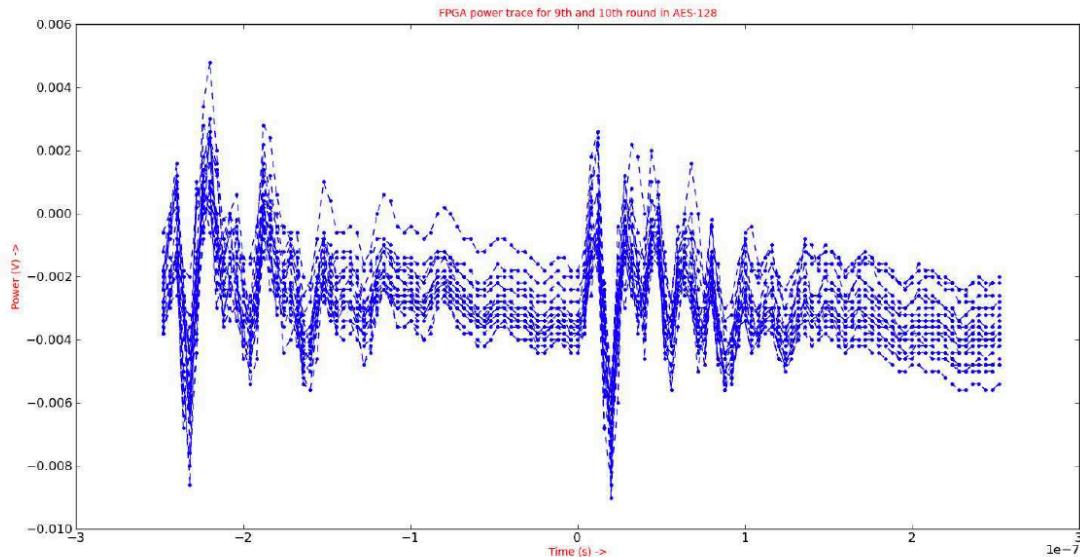
Make Some Noise

# Measurement Noise

- **Importance of high-quality measurement**
  - Noise increases the required trace count.

- These fluctuations are due to electrical noise, caused due to power supply, clock generator, conduction and radiation emissions from the components

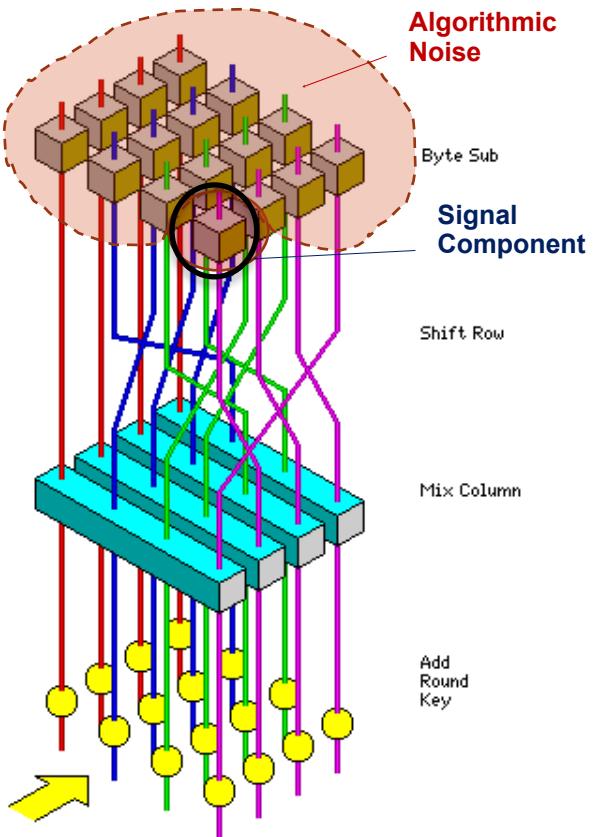
AES-128 Encryption with same plaintext and key, but resulting in different power traces.



**Simple Power Model.** Let  $t$  denote the time, and  $\mathcal{N}(t)$  be a normal distributed random variable which represents the noise components. Let  $f(g, t)$  denote the power consumption of gate  $g$  at the time  $t$ . Then a simplified power model for the power consumption is the function

$$P(t) = \sum_g f(g, t) + \mathcal{N}(t)$$

# Algorithmic Noise



- Algorithmic or switching noise occurs because of contributions of logic cells to the power consumption, which are not under attack.
- The power trace corresponds to the total power consumption of the circuit.
- However, in the attack we target only a small part (see black circle in fig.) to reveal a portion of the key.
- The power consumption from all the other parts (see blue shaded portion in fig.) form the algorithmic noise.
- This would be more in a parallel implementation, compared to a serialized implementation.

Fortunately, our statistical attack can handle both

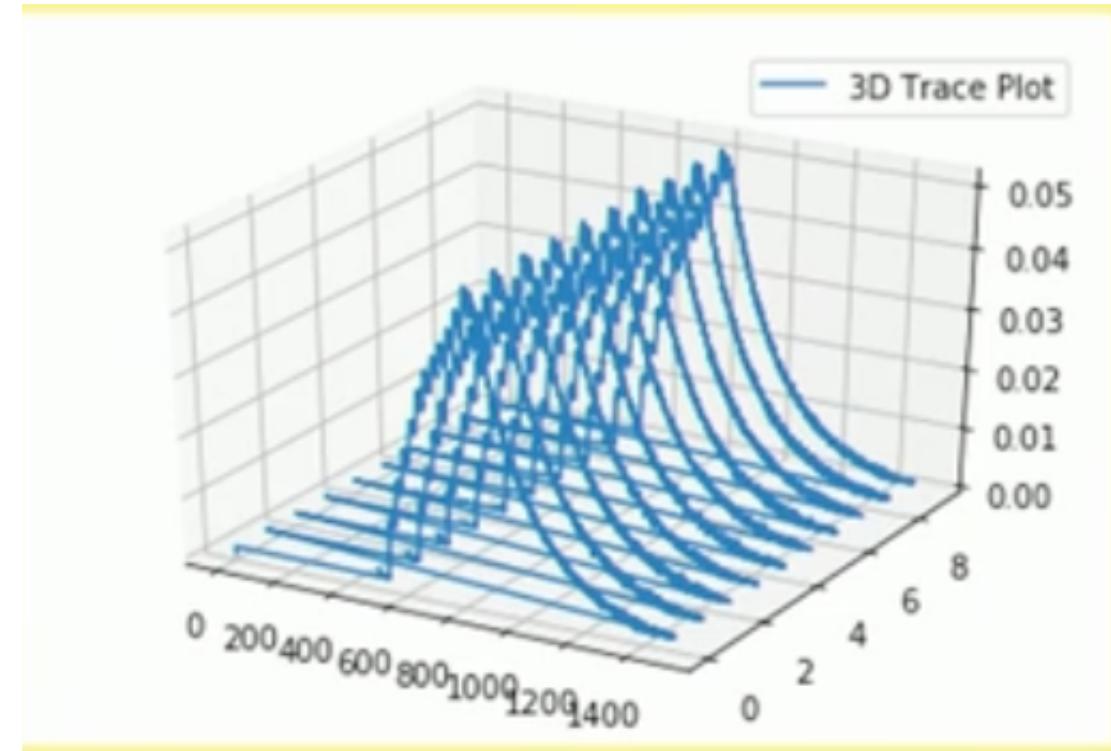
# Statistical Analysis of Power Traces

- **Fact:** For fixed operation and fixed value, the power consumption is fixed. But there will be some additive noise.
- **Fortunately, this noise is Gaussian.**
  - Let us assume that the mean is  $\mu$  and variance is  $\sigma^2$
  - Noise mainly results in the variance...
- **Recall:** The gaussian probability distribution function

$$f(x) = \frac{1}{\sigma\sqrt{2\pi}} e^{-\frac{1}{2}\frac{(x-\mu)^2}{\sigma^2}} \quad x \in \mathbb{R}$$

$$X \sim \mathcal{N}(\mu, \sigma)$$

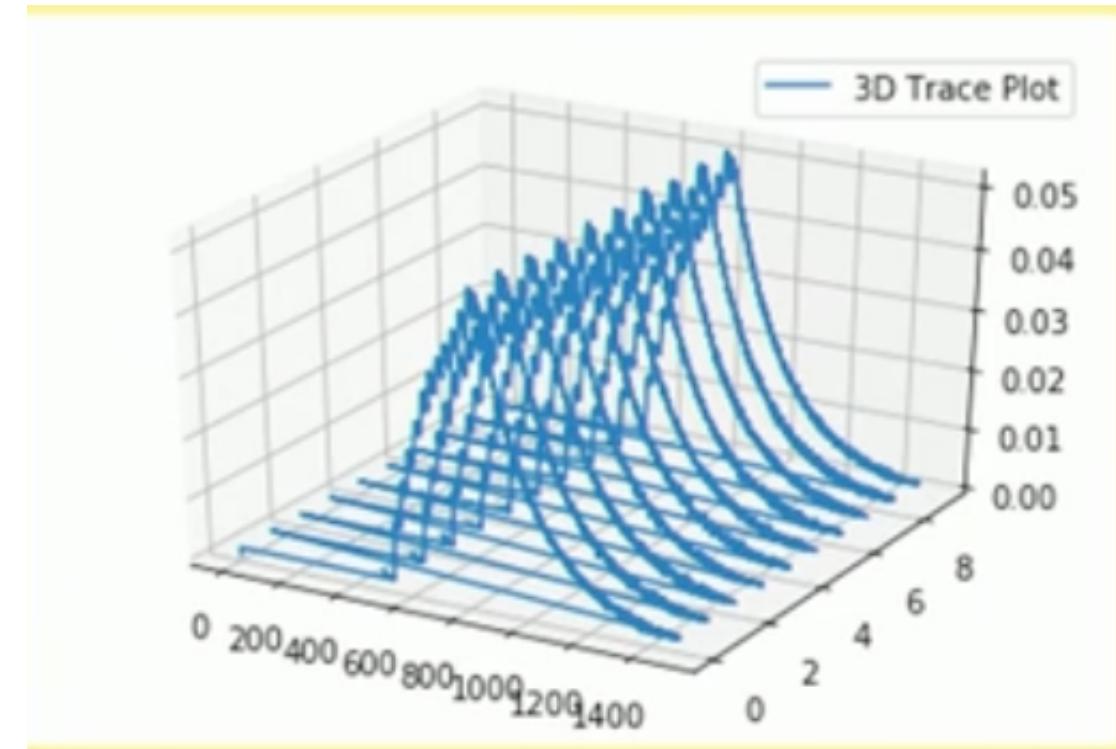
- If  $\mu = 0$ ,  $\sigma = 1$ , we call it a standard normal distribution.
- The CDF of standard normal is denoted as  $\phi(x)$ .



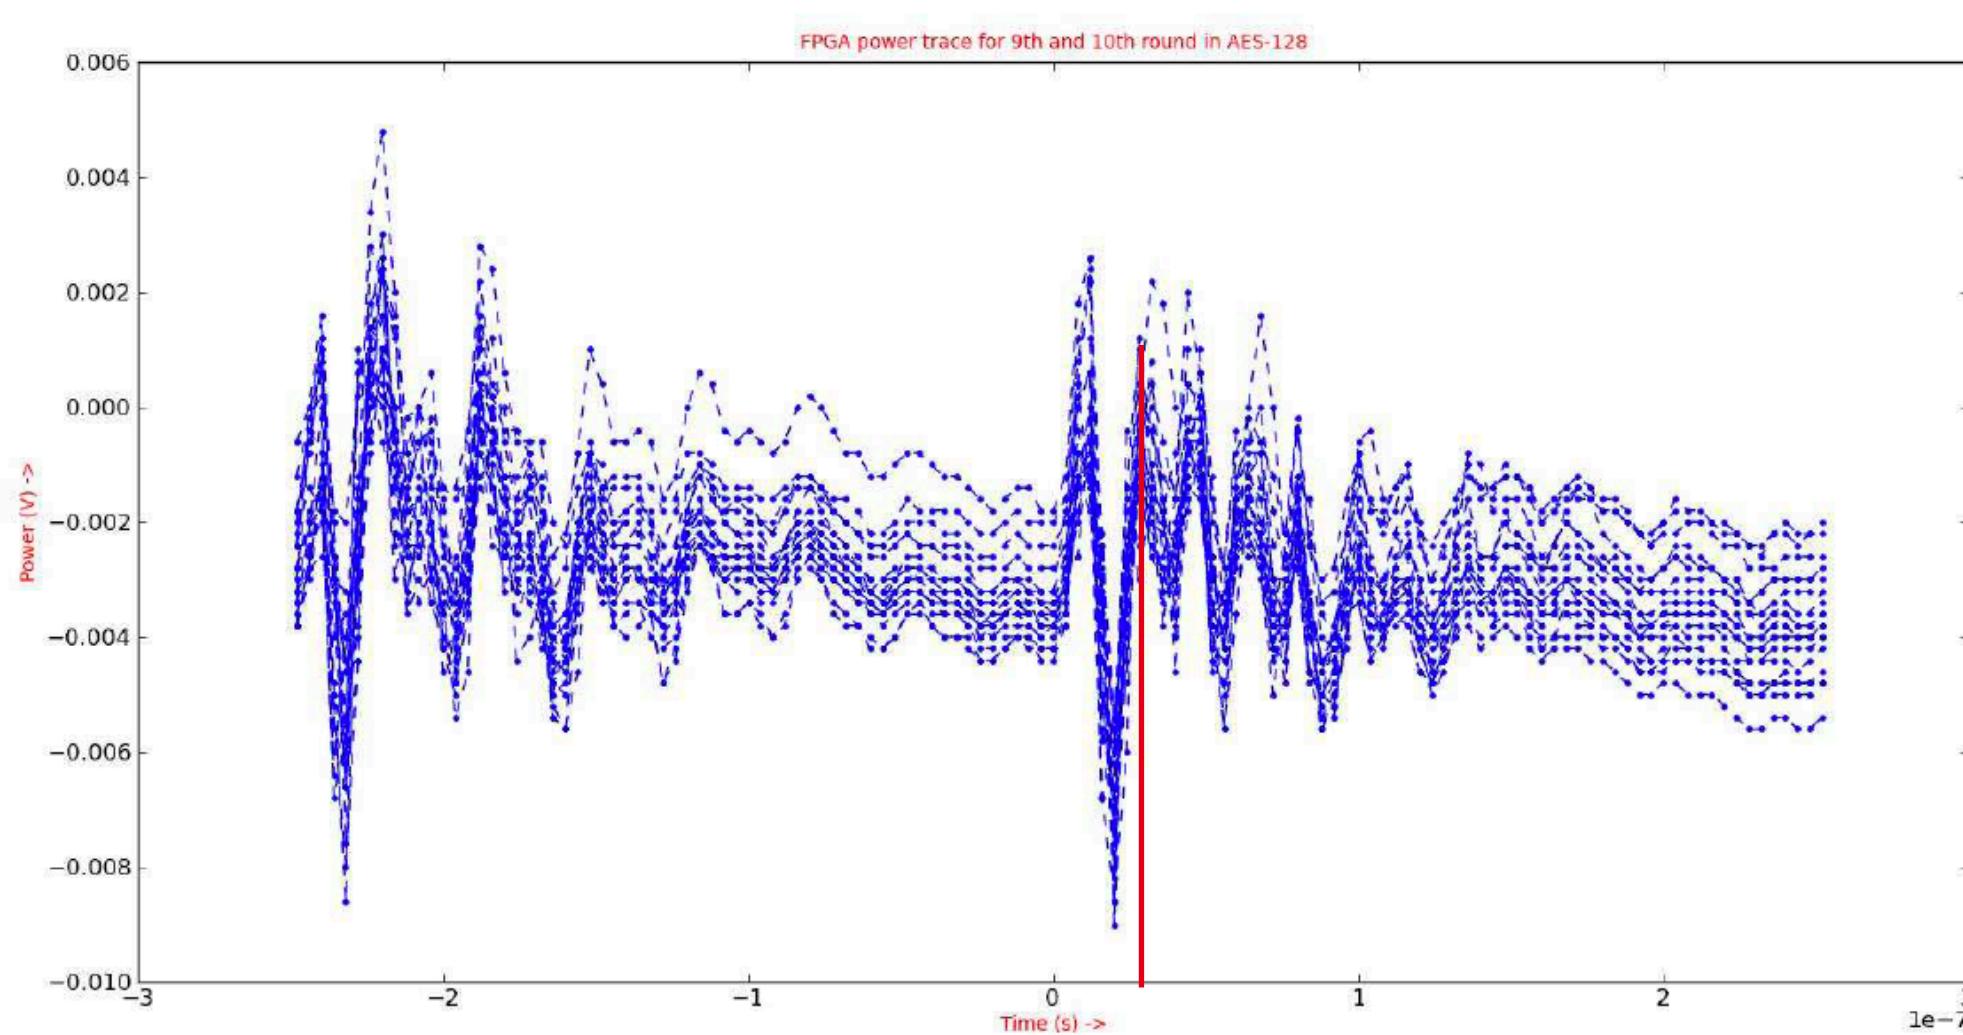
# Statistical Analysis of Power Traces

- Recall: Sampling distribution

- An *unbiased estimator* for  $\mu$  is  $\bar{x}$  — the sample mean
- An unbiased estimator for  $\sigma^2$  is  $s^2$  — the sample variance (with  $n-1$  correction factor, of course)
- **Unbiased estimator:** The expected value of the estimator is equal to the true value of the parameter; e.g.  $E(\bar{x}) = \mu$

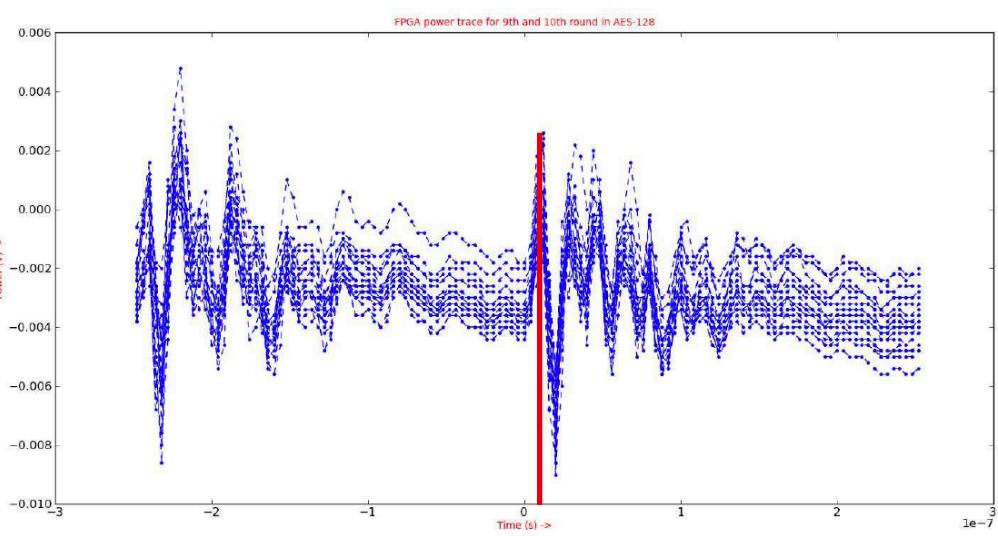


# Statistical Analysis of Power Traces



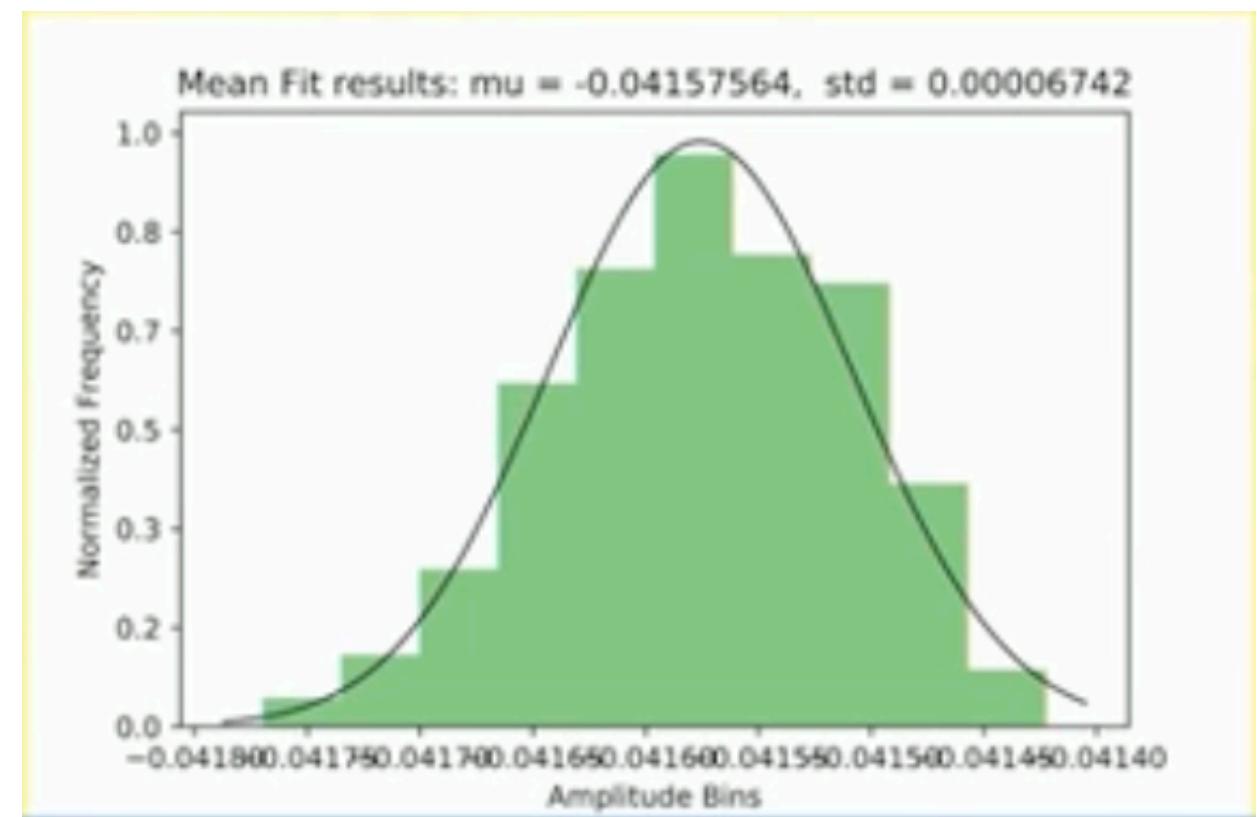
# Statistical Analysis of Power Traces

- A point in the power trace is represented as:
  - $P(t) = \sum_g f(g, t) + \mathcal{N}(\mu, \sigma, t)$
- **POI (point of interest) in the trace:** The point(s) in the trace which are the best for distinguishing the correct key from the wrong one.



**Simple Power Model.** Let  $t$  denote the time, and  $\mathcal{N}(t)$  be a normal distributed random variable which represents the noise components. Let  $f(g, t)$  denote the power consumption of gate  $g$  at the time  $t$ . Then a simplified power model for the power consumption is the function

$$P(t) = \sum_g f(g, t) + \mathcal{N}(t)$$



Histogram of a point in power trace for same pt and key